

hbr

hint for Branch (r-form)

Month: February 2011

Assembly Primer Part 4 — Hello World — ARM

On to [Assembly Primer — Part 4](#). This is where we start writing a small assembly program for the platform. In this case, I don't know the language and I don't know the ABI. Learning these from scratch ranges from interesting to tedious :)

Regarding the language (available instructions, mnemonics and assembly syntax): I'm using the [ARM Architecture Reference Manual](#) as my reference for the architecture (odd, I know). It's very long and the documentation for each instruction is extensive — which is good because there are a lot of instructions, and many of them do a lot of things at once.

Regarding the ABI (particularly things like argument passing, return values and system calls): there's the [Procedure Call Standard for the ARM Architecture](#), and there are a few other references I've found, such as the [Debian ARM EABI Port wiki page](#).

“EABI is the new “Embedded” ABI by [ARM ltd](#). EABI is actually a family of ABI's and one of the “subABIs” is GNU EABI, for Linux.”

– from Debian ARM EABI Port

System Calls

To perform a system call using the GNU EABI:

- put the system call number in r7
- put the arguments in r0-r6 (64bit arguments must be aligned to an even numbered register i.e. in r0+r1, r2+r3, or r4+r5)
- issue the **Supervisor Call** instruction with a zero operand — **svc #0**

(**Supervisor Call** was previously named **Software Interrupt** — **swi**)

Just Exit

Based on the above, it's not difficult to reimplement **JustExit.s** ([original](#)) for ARM.

```
.text

.globl _start

_start:
    mov r7, #1
    mov r0, #0
    svc #0
```

mov here is **Move (Immediate)** which puts the #-prefixed literal into the named register.

Hello World

Likewise, the conversion of **HelloWorldProgram.s** ([original](#)) is not difficult:

```
.data

HelloWorldString:
    .ascii "Hello World\n"

.text
```

```
.globl _start

_start:
    # Load all the arguments for write ()

    mov r7, #4
    mov r0, #1
    ldr r1,=HelloWorldString
    mov r2, #12
    svc #0

    # Need to exit the program

    mov r7, #1
    mov r0, #0
    svc #0
```

This includes the **load register** pseudo-instruction, **ldr** — the compiler stores the address of **HelloWorldString** into the literal pool, a portion of memory located in the program text, and the 32bit address is loaded from the literal pool ([more details](#)).

When compiling a similar C program with **-mcpu=cortex-a8**, I notice that the compiler generates **Move (immediate)** and **Move Top** — **movw** and **movt** — instructions to load the address directly from the instruction stream, which is presumably more efficient on that architecture.

February 14, 2011 / general, programming / ARM, Assembly / 3 Comments

Assembly Primer Parts 1, 2 and 3 — ARM

I had started a [series of posts on assembly programming](#) for the Cell BE PPU and SPU, based on the [assembly primer video series from securitytube.net](#). I have recently acquired a Nokia N900, and so thought I might take the opportunity to continue the series with a look at the ARM processor as well.

Wikipedia lists the N900's processor as a Texas Instruments OMAP3430, 600MHz ARMv7 Cortex-A8. I'm not at all familiar with the processor family, so I'll be attempting to find out what all of this means as I go :P

I've set up a cross compiler on my desktop machine using Gentoo's neat crossdev tool (built using `crossdev -t arm-linux-gnueabi`). The toolchain builds a functional Hello, World!

(I note that scratchbox appears to be the standard tool/environment used to build apps for Maemo — I may take a closer look at that at a later date)

I have whatever the latest public 'stable' Maemo 5 release is on the N900 (PR 1.3, I think), with an `apt-get install openssh gdb` — thus far, enough to "debug" a functional Hello, World!

What follows are some details of the Cortex-A8 architecture present in the N900, particularly in how it differs from IA32, as presented in the videos [Part 1 — System Organisation](#), [Part 2 — Virtual Memory Organization](#) and [Part 3 — GDB Usage Primer](#). I've packed them all into this post because gdb usage and Linux system usage are largely the same on ARM as they are on PPC and IA32.

Most of the following information comes from the [ARM Architecture Reference Manual](#).

(The number of possible configurations of ARM hardware makes it interesting at times to work out exactly which features are present in my particular processor. From what I can tell, the N900's Cortex-A8 is ARMv7-A and includes VFPv3 half, single and double precision float support, and NEON (aka Advanced SIMD). I expect I'll find out more when I actually start to try and program the thing. As to which gcc -march, -mcpu or -mfpu options are most correct for the N900 — I have no idea.)

1. Registers

Integer

There are sixteen 32bit ARM core registers, R0 to R15, where R0–R12 are for general use. R13 contains the stack pointer (SP), R14 is the link register (LR), and R15 is the program counter (PC).

The current program status register (CSPR) contains various status and control bits.

VFPv3 (Floating point) & NEON (Advanced SIMD)

There are thirty two doubleword (64bit) registers, that can be referenced in a number of ways.

NEON instructions can access these as thirty two doubleword registers (D0–D31) or as sixteen quadword registers (Q0–Q15), able to be used interchangeably.

VFP instructions can view the same registers as 32 doubleword registers (again, D0–D31) or as 32 single word registers (S0–S31). The single word view is packed into the first 16 doubleword registers.

Something like this pic (click to embiggen):

ARMv7 Advanced SIMD (NEON) and VFPv3 extension registers																																S:32bit D:64bit Q:128bit							
VFP	S0	S1	S2	S3	S4	S5	S6	S7	S8	S9	S10	S11	S12	S13	S14	S15	S16	S17	S18	S19	S20	S21	S22	S23	S24	S25	S26	S27	S28	S29	S30	S31							
VFP+NEON	D0	D1	D2	D3	D4	D5	D6	D7	D8	D9	D10	D11	D12	D13	D14	D15																							
NEON	Q0				Q1				Q2				Q3				Q4				Q5				Q6				Q7										
VFP+NEON	D16	D17	D18	D19	D20	D21	D22	D23	D24	D25	D26	D27	D28	D29	D30	D31																							
NEON	Q8				Q9				Q10				Q11				Q12				Q13				Q14				Q15										

VFP in this core (apparently) supports single and double precision floating point data types and arithmetic, as well as half precision (possibly in two different formats...).

NEON instructions support accessing values in extension registers as

- 8, 16, 32 or 64bit integer, signed or unsigned,
- 16 or 32bit floating point values, and
- 8 or 16bit polynomial values.

There's also a floating point status and control register (FPSCR).

2. Virtual Memory Organisation

On this platform, program text appears to be loaded at **0x8000**.

After an `echo 0 > /proc/sys/kernel/randomize_va_space`, the top of the stack appears to be **0xbfb00000**.

3. SimpleDemo

Compared to the video, there are only a couple of small differences when running SimpleDemo in gdb on ARM.

Obviously, the disassembly is not the same as for IA32. Rather than the **call** instructions noted in the video, you'll see **bl** (Branch with Link) for the various functions called.

Where the return address is stored on the stack for IA32, the link register (**lr** in **info registers** output) stores the return address for the current function, although **lr** will be pushed to the stack before another function is called.

(From a cursory googling, it seems that to correctly displaying all VFP/NEON registers requires gdb-7.2 — I'm running the 6.8-based build from the maemo repo. crossdev will build me a gdb I can run on my desktop PC — **crossdev -t arm-linux-gnueabi -ex-gdb** — but I believe I still need to build a newer **gdbserver** to run on the N900.)

Other assembly primer notes are linked [here](#).

February 12, 2011 / general, programming / ARM, Assembly / [Leave a comment](#)

hbr / Proudly powered by WordPress