

ARM Virtualization: CPU & MMU Issues


Prashanth Bungale, Sr. Member of Technical Staff

Overview

- **Virtualizability and Sensitive Instructions**
- **ARM CPU State**
- **Sensitive Instructions in ARM**
 - Dealing with Sensitive Instructions
 - In place binary translation
- **MMU Virtualization: Why & How?**
- **Comparison of ARM & x86 Virtualization**

Virtualizability and Sensitive Instructions

- **Defined in the context of a particular virtualization technique**
- **Example: Trap and Emulate Model**
 - Let VM execute most of its instructions directly on the h/w
 - Except for some sensitive instructions that trap into the VMM and are emulated
 - Sensitive instructions are those that interfere with:
 - Correct emulation of the VM, or
 - Correct functioning of the VMM
 - E.g. “Halt Machine” instruction



Virtualizability and Sensitive Instructions

■ Sensitive Instructions as defined by Goldberg [1]

- Mode Referencing Instructions
- Sensitive Register/Memory Access Instructions
- Storage Protection System Referencing Instructions
- All I/O Instructions

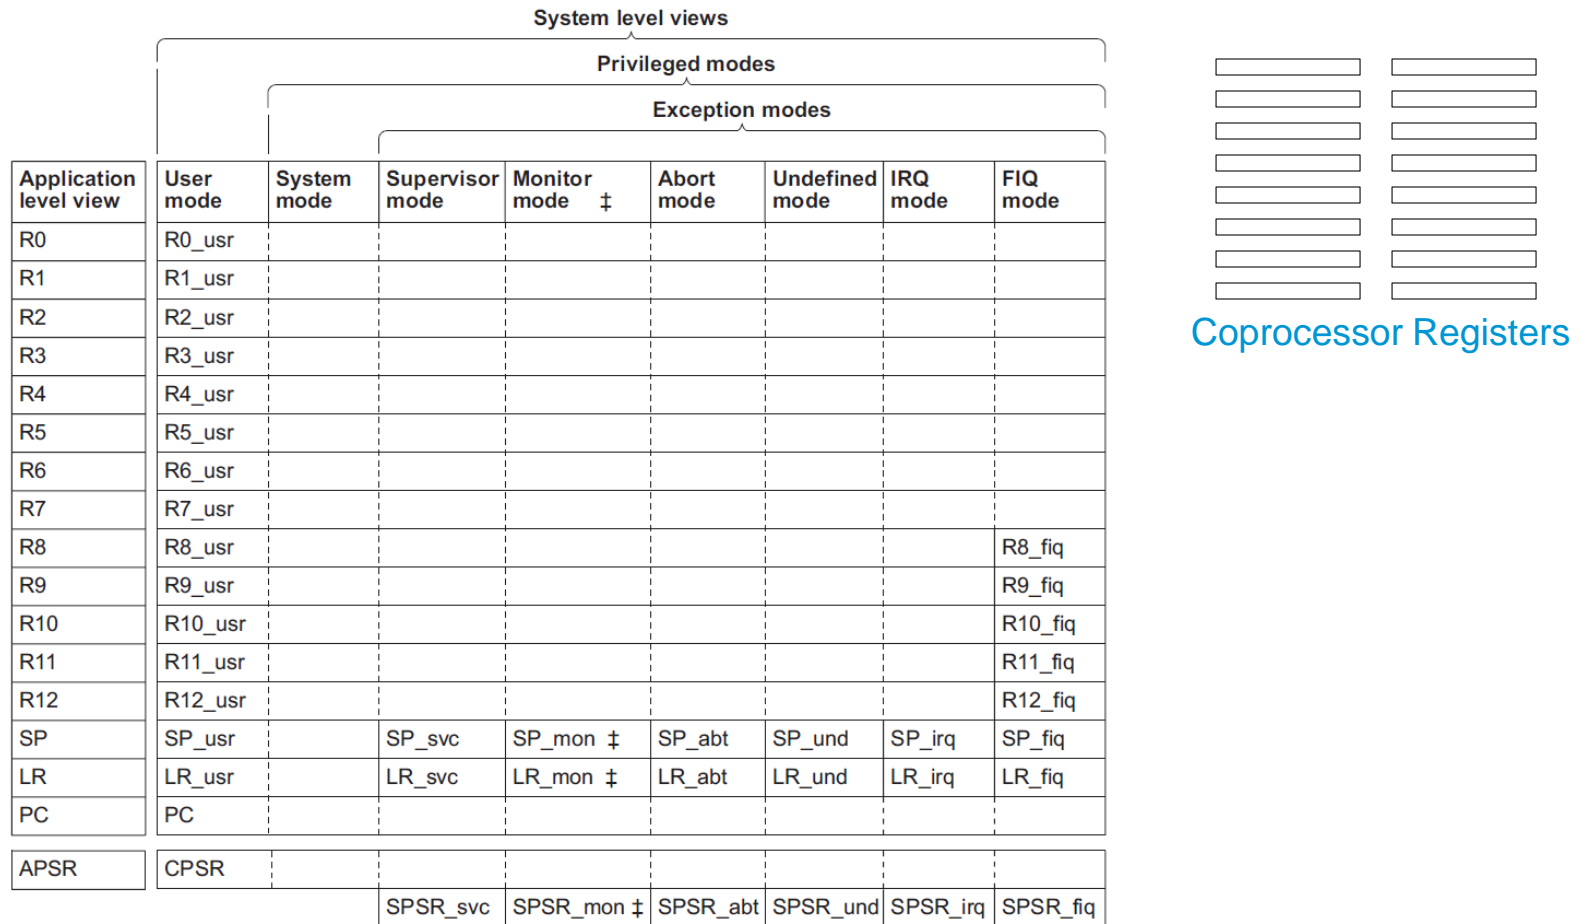
■ Popek and Goldberg's Theorem about *strict virtualizability* [2]

- *For any conventional third generation computer, a virtual machine monitor may be constructed if the set of sensitive instructions for that computer is a subset of the set of privileged instructions.*

[1] Goldberg. Architectural Principles for Virtual Computer Systems. Ph.D. thesis, Harvard University, 1972.

[2] Popek and Goldberg, Formal requirements for virtualizable third generation architectures. In SOSP '73

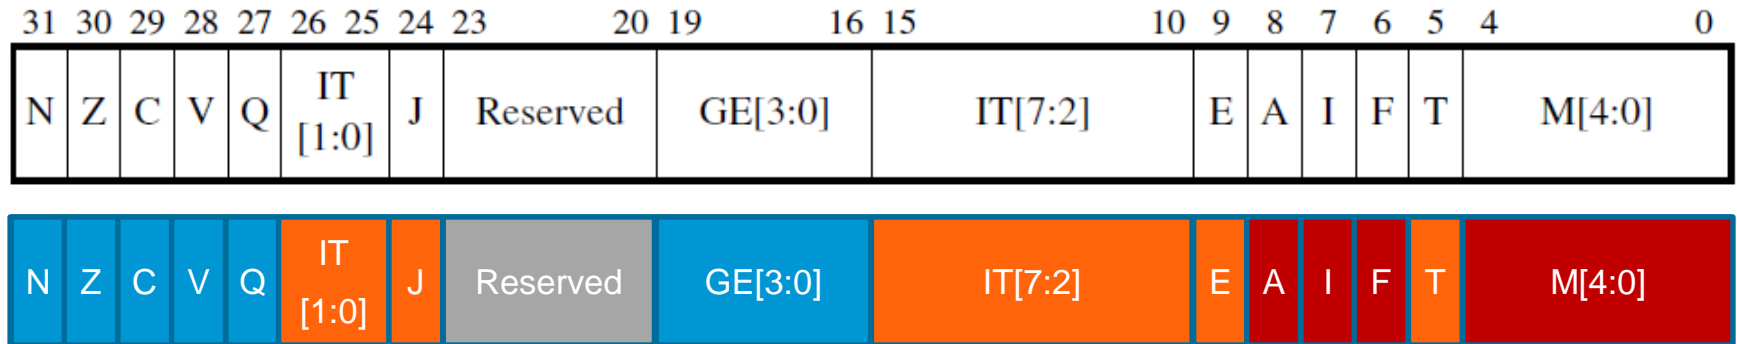
ARM CPU State



‡ Monitor mode, and the associated banked registers, are implemented only as part of the Security Extensions

Figure B1-1 Organization of general-purpose registers and Program Status Registers


CPSR: Current Program Status Register



Interrupt
Masks

Current
Processor Mode

Execution State Registers:



- Privileged-only Access Registers
- Execution State Registers
- Condition Flags

Sensitive Instructions on ARM

- **Coprocessor Access Instructions**
 - MRC / MCR / CDP / LDC / STC
- **SIMD/VFP System Register Access Instructions**
 - VMRS / VMSR
- **TrustZone Secure State Entry Instructions**
 - SMC
- **Memory-Mapped I/O Access Instructions**
 - Load/Store instructions from/into memory-mapped I/O locations
- **Direct (Explicit/Implicit) CPSR Access Instructions**
 - MRS / MSR / CPS / SRS / RFE / LDM(Exc. Return) / DPSPC
- **Indirect CPSR Access Instructions**
 - LDRT / STRT – Load/Store Unprivileged (“As User”)
- **Banked Register Access Instructions**
 - LDM/STM(User mode registers)

Dealing with Sensitive Instructions


■ Hardware Techniques

- Privileged Instruction Semantics dictated by ISA - 1, 2, 3
- MMU-enforced traps (e.g., page fault) - 4
- Tracing/debug support (e.g., bkpt)
- Hardware-assisted Virtualization (e.g., extra privileged mode)

Dealing with Sensitive Instructions

■ Software Techniques

- Interpretation / Full Emulation
- Binary Translation
 - Cached Translation →
 - In-Place Translation
- Para-Virtualization
 - Shallow Para-Virtualization: replace sensitive instructions
 - Deep Para-Virtualization: replace sensitive subsystems (e.g., process model, page-table management, etc.)
- Binary Patching / Pre-Virtualization




In-place binary translation

- **ARM has a fixed instruction size**

- 32-bit in ARM mode and 16-bit in Thumb mode

- **Perform binary translation in-place**

- Instead of in a separate cache
- Follow control-flow
 - Translate basic block (if not already translated) at the current PC
 - Ensure interposition at end of translated sequence
 - All writes (but not reads) to PC now become dangerous instructions
- Replace dangerous (i.e., sensitive but unprivileged) instructions 1-1 with hypercalls to force trap and emulate




- **Guest transparency issues**

1-1 Hypercalls

■ Replace dangerous instructions with 1-1 Hypercalls

- Use trap instruction to issue hypercall
- Encode hypercall type & original instruction bits in hypercall *hint*
- Example:



■ Trap and Emulate Semantics

- Upon trapping into the monitor, decode the hypercall type and the original instruction bits, and emulate instruction semantics

MMU Virtualization: Why?

■ Translation Virtualization

- $VA \rightarrow PA$; $PA \rightarrow MA$

■ Protection Virtualization

- Access permissions, domains
- Multiplex available hardware MMU protection support among:
 - Protecting the hypervisor from the guests
 - Protecting the guest kernel from guest user

■ Cross-architectural versions (e.g., ARMv5 on ARMv6)

- Dealing with differences in cache, TLB architectures, page-table layout, etc.

MMU Virtualization: How?

■ Shadow PT

- Intercept guest MMU events of interest
 - Data/Prefetch Aborts, TTBR deltas, PT deltas, TLB ops
- Maintain (lazily) hypervisor-controlled, trusted shadow PT
- Options:
 - TLB Coherency-based Shadow PTs / Cached Shadow PTs
 - In-place Shadow PTs

■ Para-Virtualized trusted guest PT

- Highly intrusive to guest MMU software

■ Hardware virtualization support

- Nested / 2-stage Page Tables: VA->PA; PA->MA

Comparison of ARM vs. x86 Virtualizability

■ Sensitive Instructions

Type of Sensitive Instructions	Violating Goldberg's Requirement #	X86 [3]	ARM
Sensitive Register Access	3B	SGDT, SIDT, SLDT, SMSW, PUSHF/POPF	-
Protection System References	3C	LAR, LSL, VERR, VERW, PUSH/POP, CALL, JMP, INT n, RET, STR, MOVE	LDM/STM (user regs), LDRT/STRT ("As User")
Both	3B & 3C	-	MRS, MSR, CPS, SRS, RFE, DPSPC, LDM (exc. return)

[3] John Scott Robin and Cynthia Irvine, Analysis of the Intel Pentium's Ability to Support a Secure Virtual Machine Monitor, USENIX Security Symposium, 2000.

Comparison of ARM vs. x86 Virtualizability

■ Ring compression – protection mechanisms

- x86: Segmentation + Paging
- ARM: Paging (+ domains?)

■ Instruction execution vs. Data Read/Write protection

- x86: CS for instruction fetch vs. DS/other for data access
- ARM: No explicit distinction b/w execute and read protection

■ Cache architecture

- x86: Largely transparent; PIPT across all versions
- ARM: Exposes a lot of the cache architecture; VIVT/VIPT/PIPT

Comparison of ARM vs. x86 Virtualizability

■ Instruction size

- x86: Variable
- ARM: Fixed -> enables more in-place patching mechanisms

■ I/O

- x86: I/O instructions + memory-mapped I/O
- ARM: Only memory-mapped I/O