

# ARM Virtualization: Performance and Architectural Implications

Christoffer Dall, Shih-Wei Li, Jin Tack Lim,  
Jason Nieh, and Georgios Koloventzos



# ARM<sup>®</sup>



ARM Servers




ARM Network Equipment



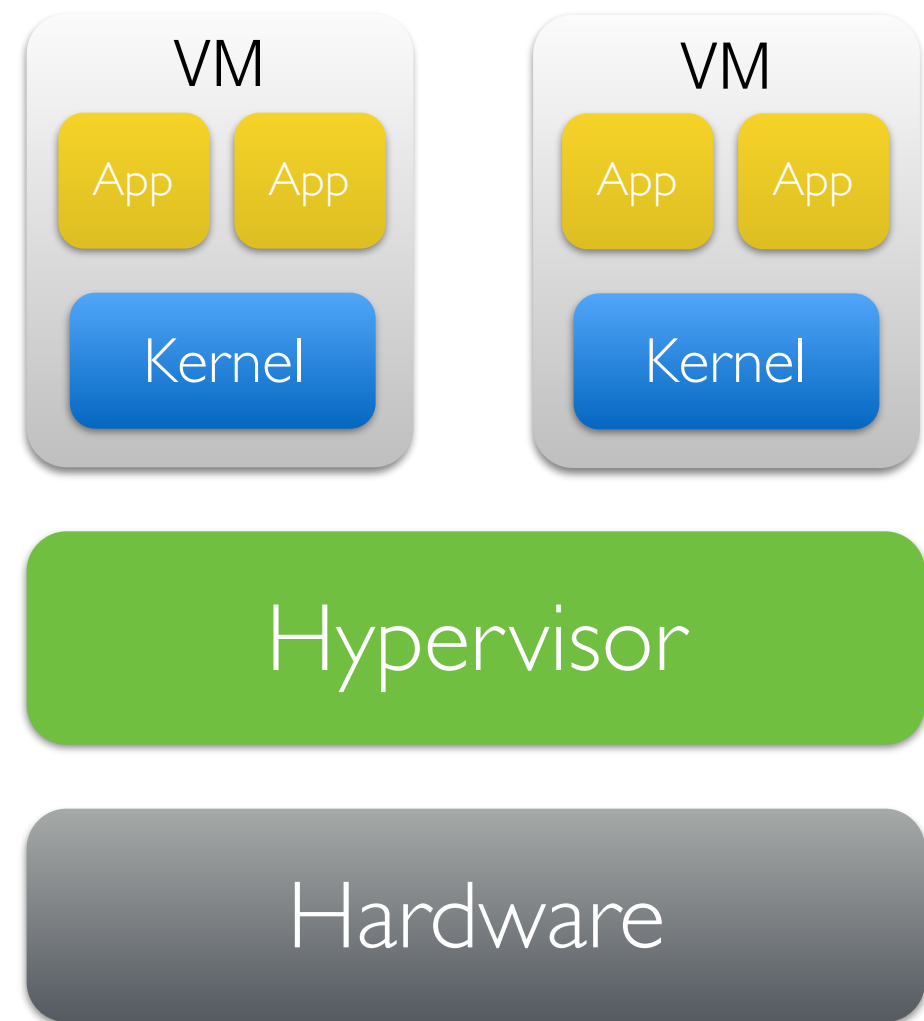
Virtualization

# Virtualization

Native




Virtual Machines



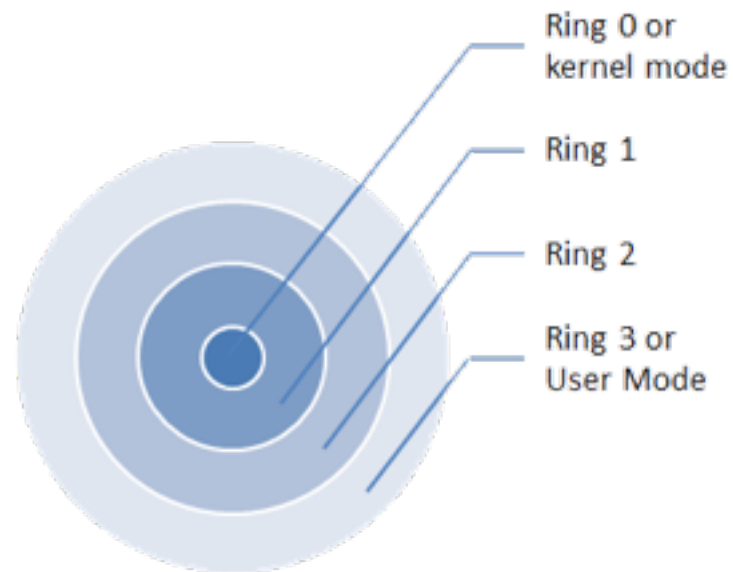
# ARM Hardware Virtualization Support

**ARM<sup>®</sup>**  
**Virtualization Extensions**

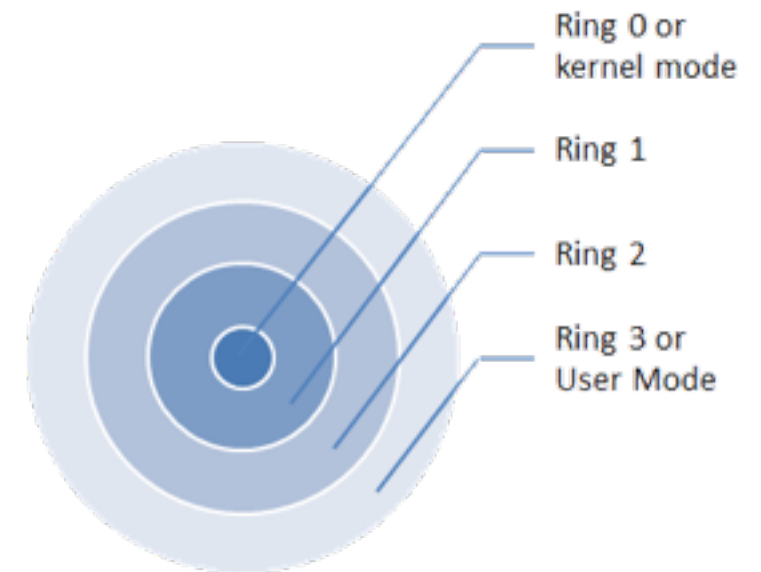


# x86

Root (Hypervisor)



Non-Root (VM)

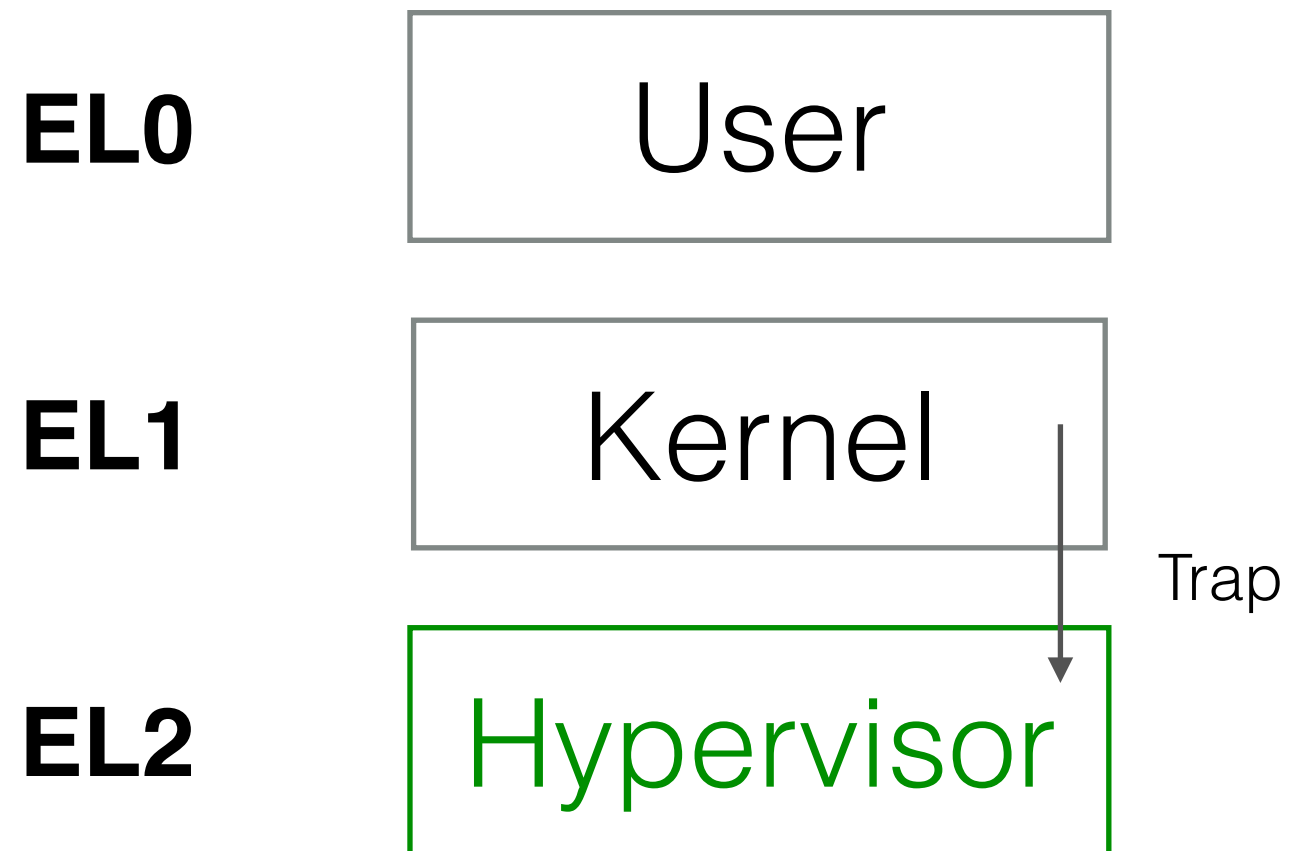


VM Exit

Save/Restore state to  
VMCS

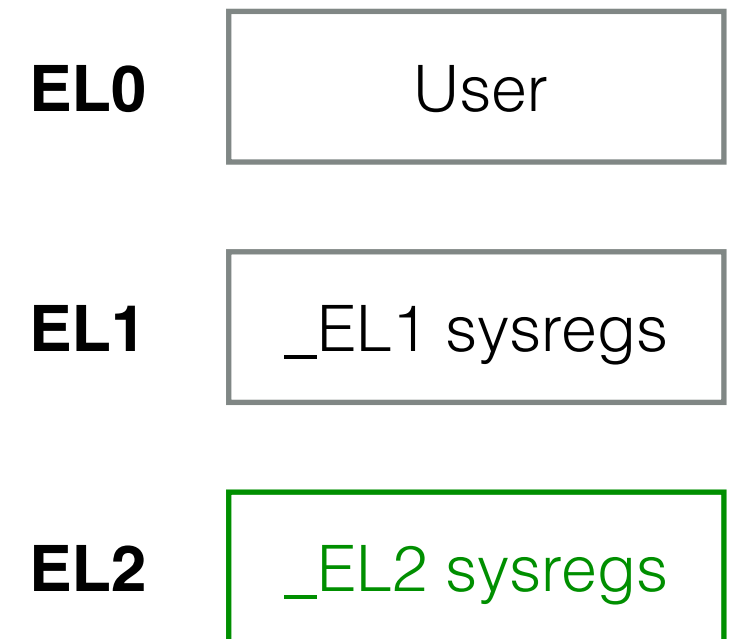


# ARM Virtualization Extensions



# EL2

- Controlled by EL2 system registers
- Limited to support hypervisors, not OS kernels



# ARM Virtualization Extensions

## Design Choices

1. Clear hierarchy from user to kernel to hypervisor

**EL0**

User

**EL1**

Kernel

2. Reduced complexity

**EL2**

Hypervisor



# ARM Virtualization Performance ?

# Measurement Study

- Micro-benchmarks: low-level hypervisor operations
- Macro-benchmarks: application workloads

# Hardware Setup

## ARM Hardware

- HP Moonshot m400
- 64-bit ARMv8-A
- 2.4 GHz APM Atlas CPU
- 8-way SMP
- 64 GB RAM (capped at 16 GB)
- 10 GB Ethernet

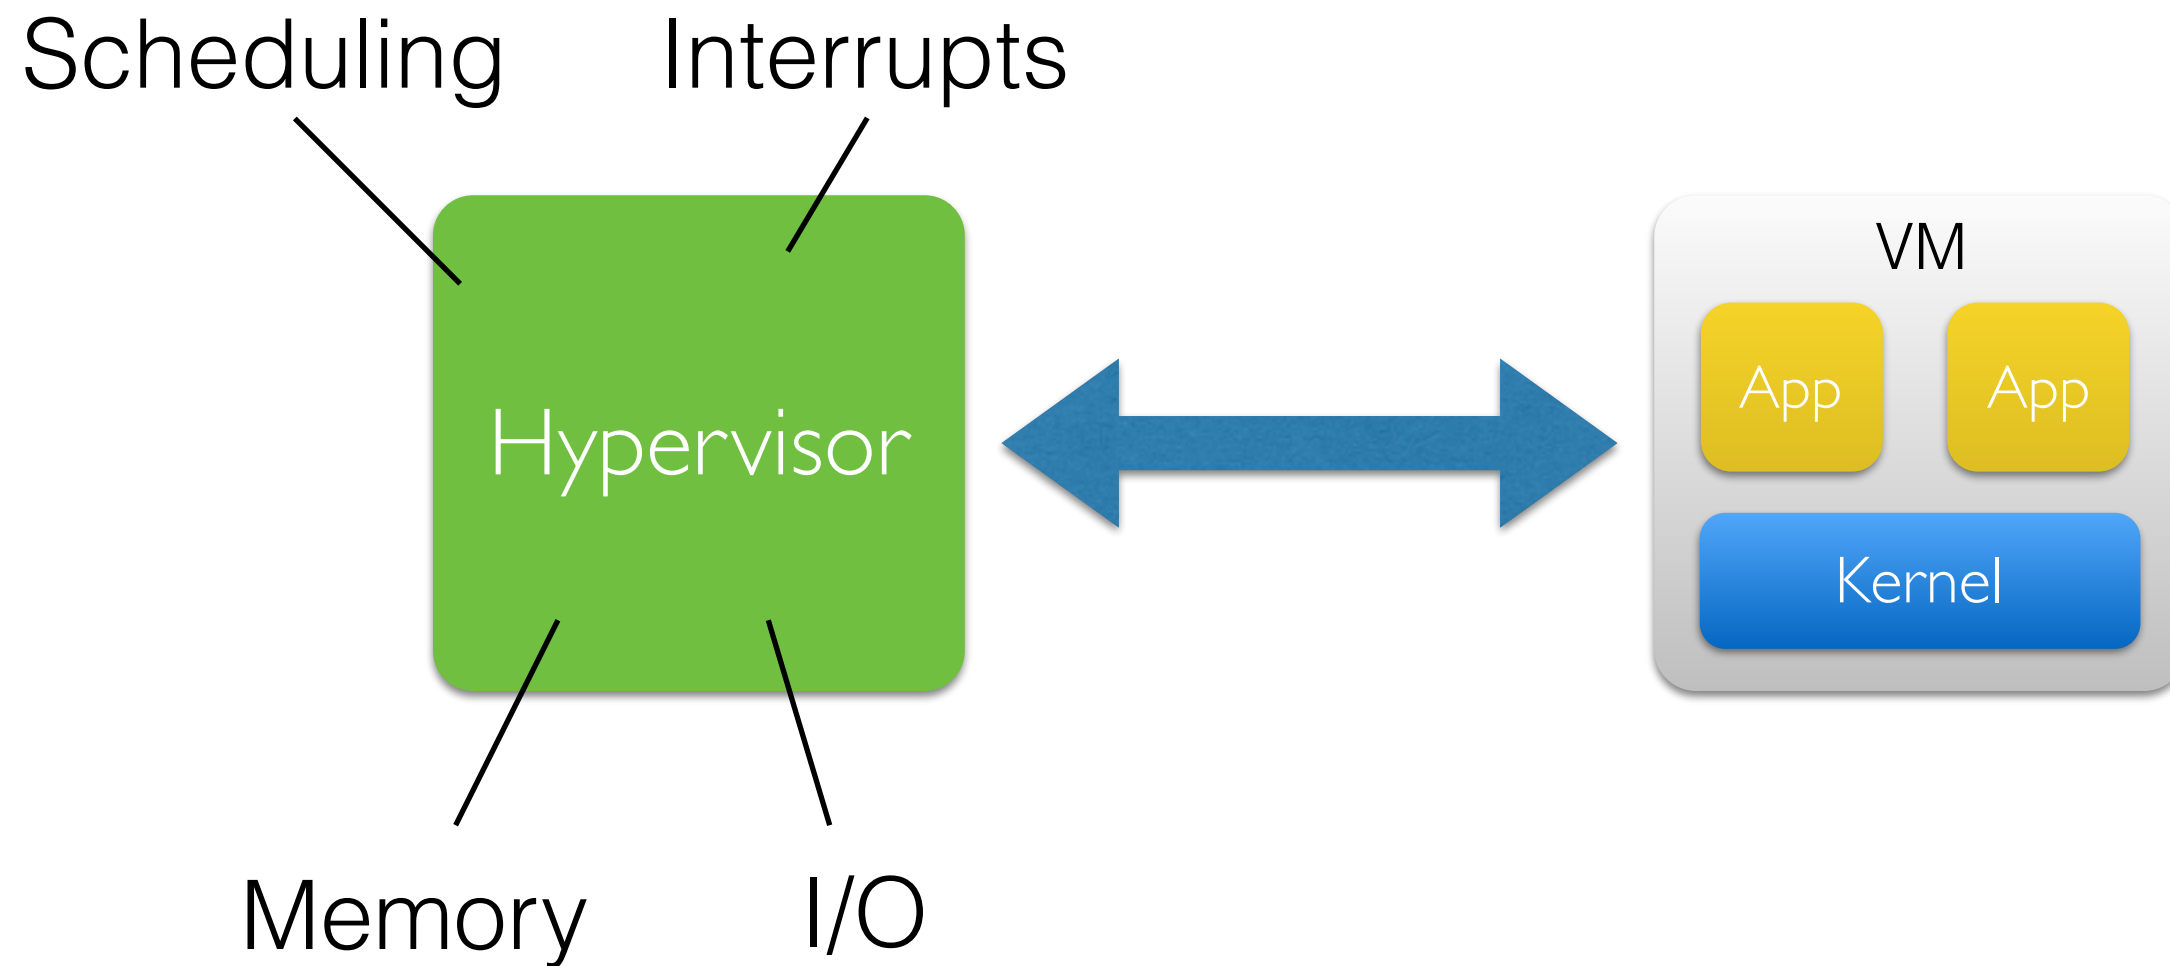
## x86 Hardware

- Dell PowerEdge r320
- 64-bit x86\_x64
- 2.1 GHz Intel Xeon ES-2450
- 8-way SMP
- 16 GB RAM
- 10 GB Ethernet


# Software Setup



# VM-to-Hypervisor Transitions



# No-Op Hypercall






# Micro Results

CPU Clock Cycles	ARM		x86	
	KVM	Xen	KVM	Xen
Hypercall	6,500	376	1,300	1,228

1: ARM can be either much faster or slower than x86

# x86

Non-Root (VM)



Root (Hypervisor)

# ARM

**EL0**

User

**EL1**

Kernel

**EL2**

Hypervisor

Trap



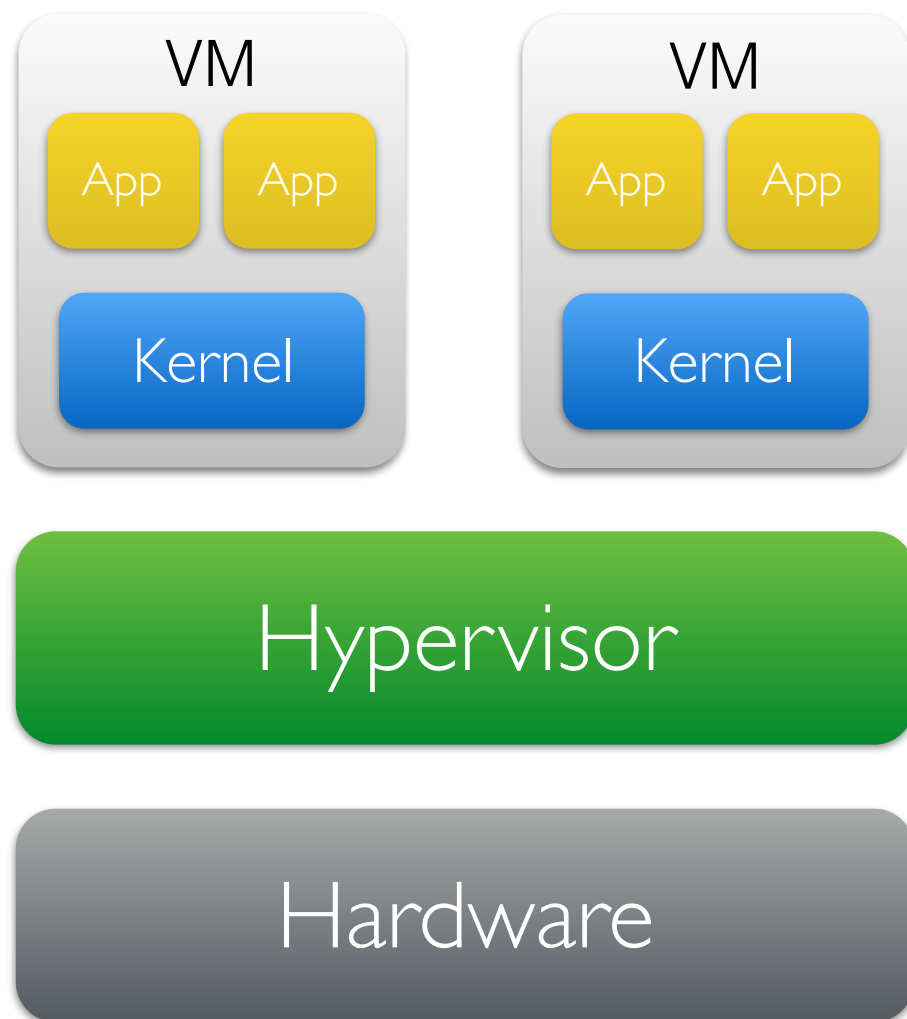
# Micro Results

CPU Clock Cycles	ARM		x86	
	KVM	Xen	KVM	Xen
Hypercall	6,500	376	1,300	1,228

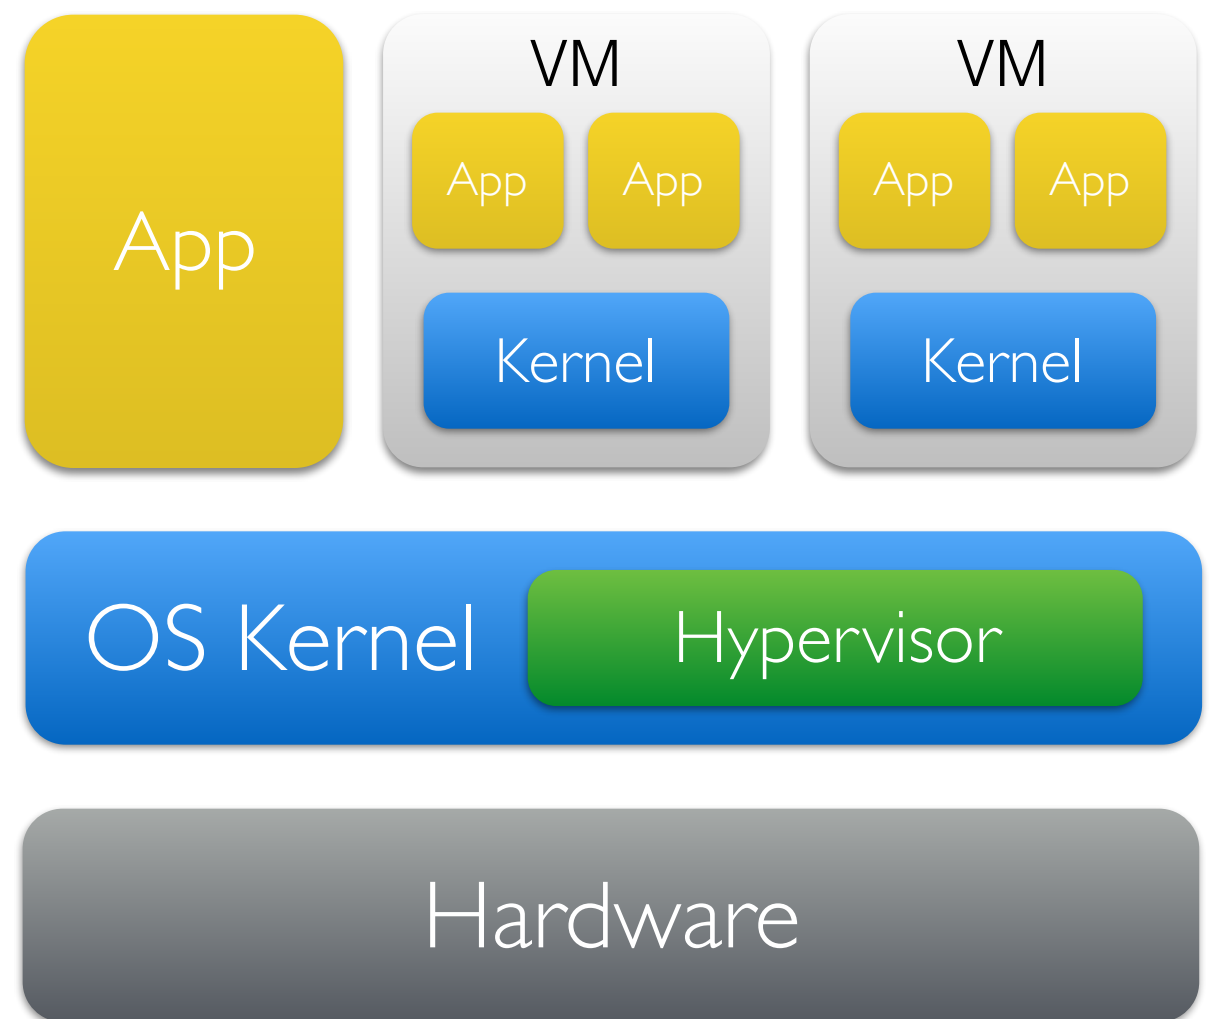
- 1: ARM can be either much faster or slower than x86  
-> x86 VM Exit more complicated than ARM Trap
- 2: KVM is much slower than Xen on ARM

# Hypervisor Design

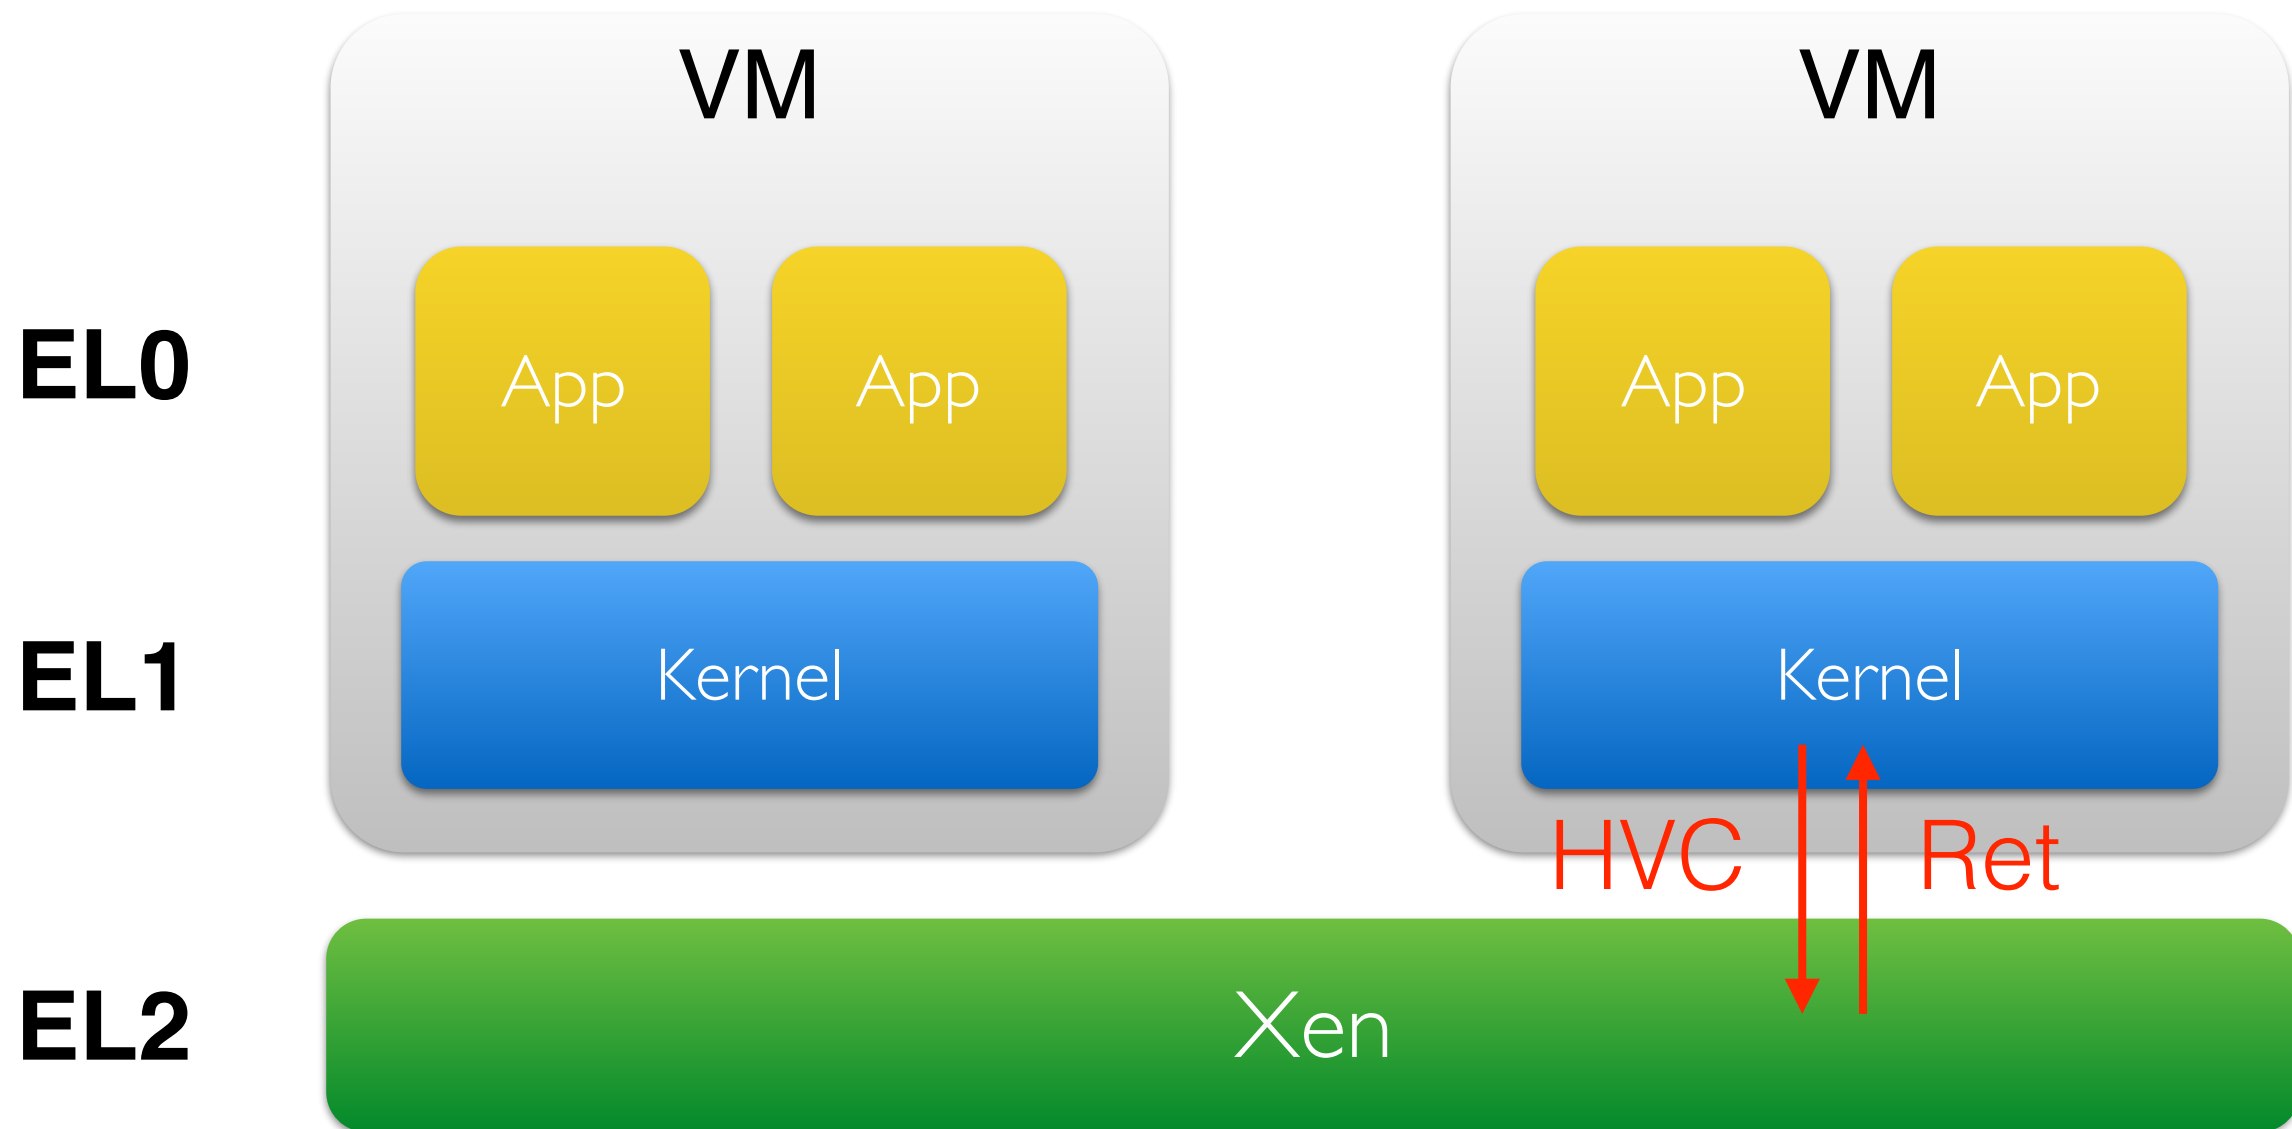
Type 1 (Bare-Metal)



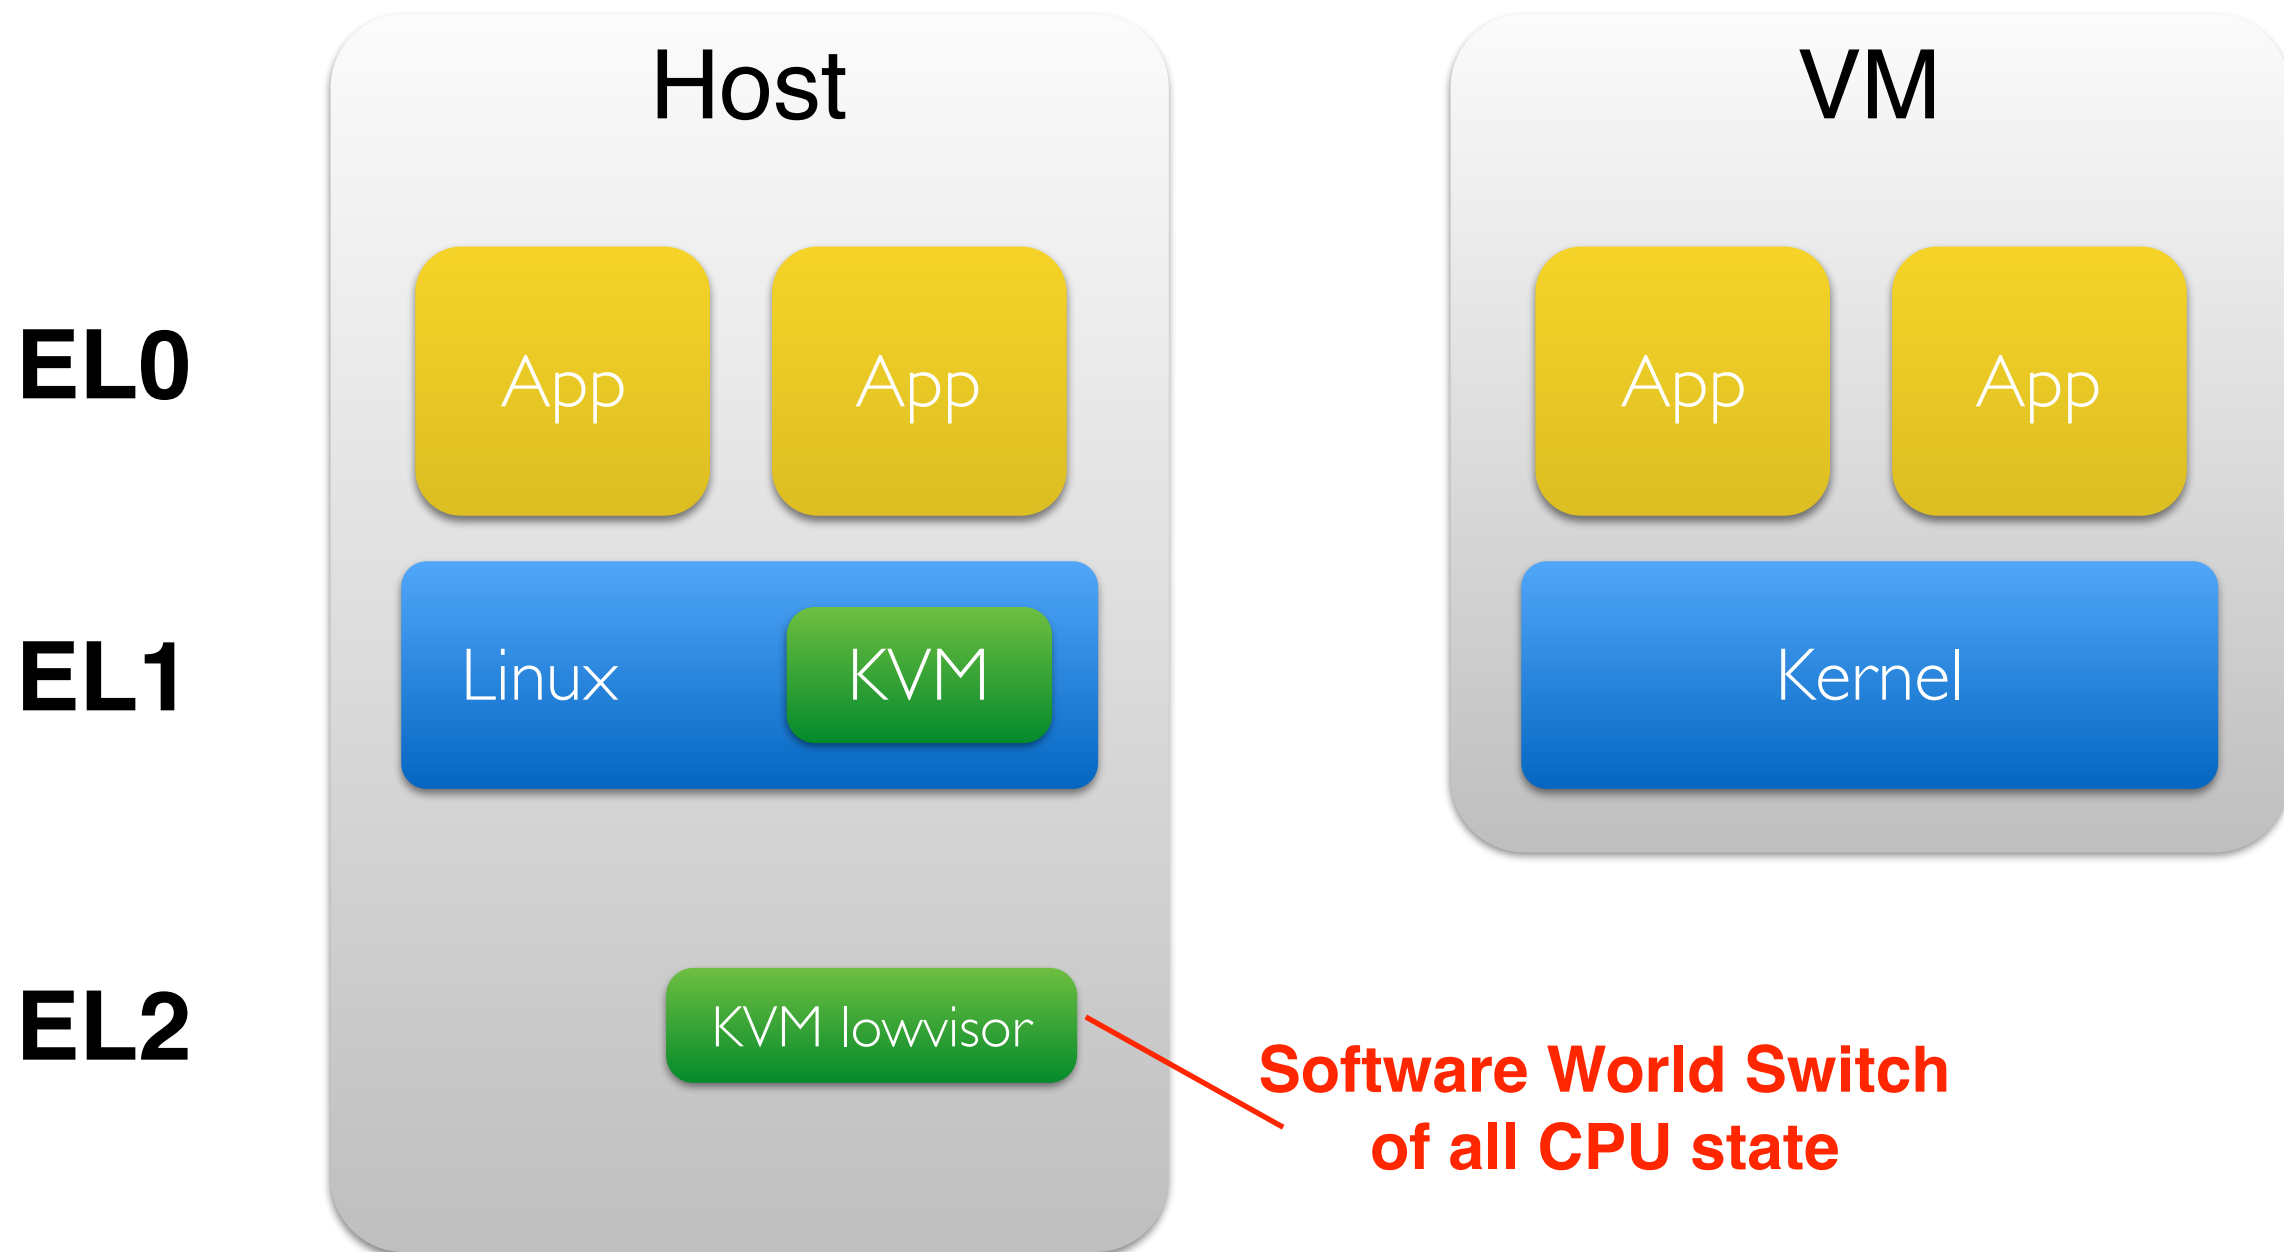
Type 2 (Hosted)



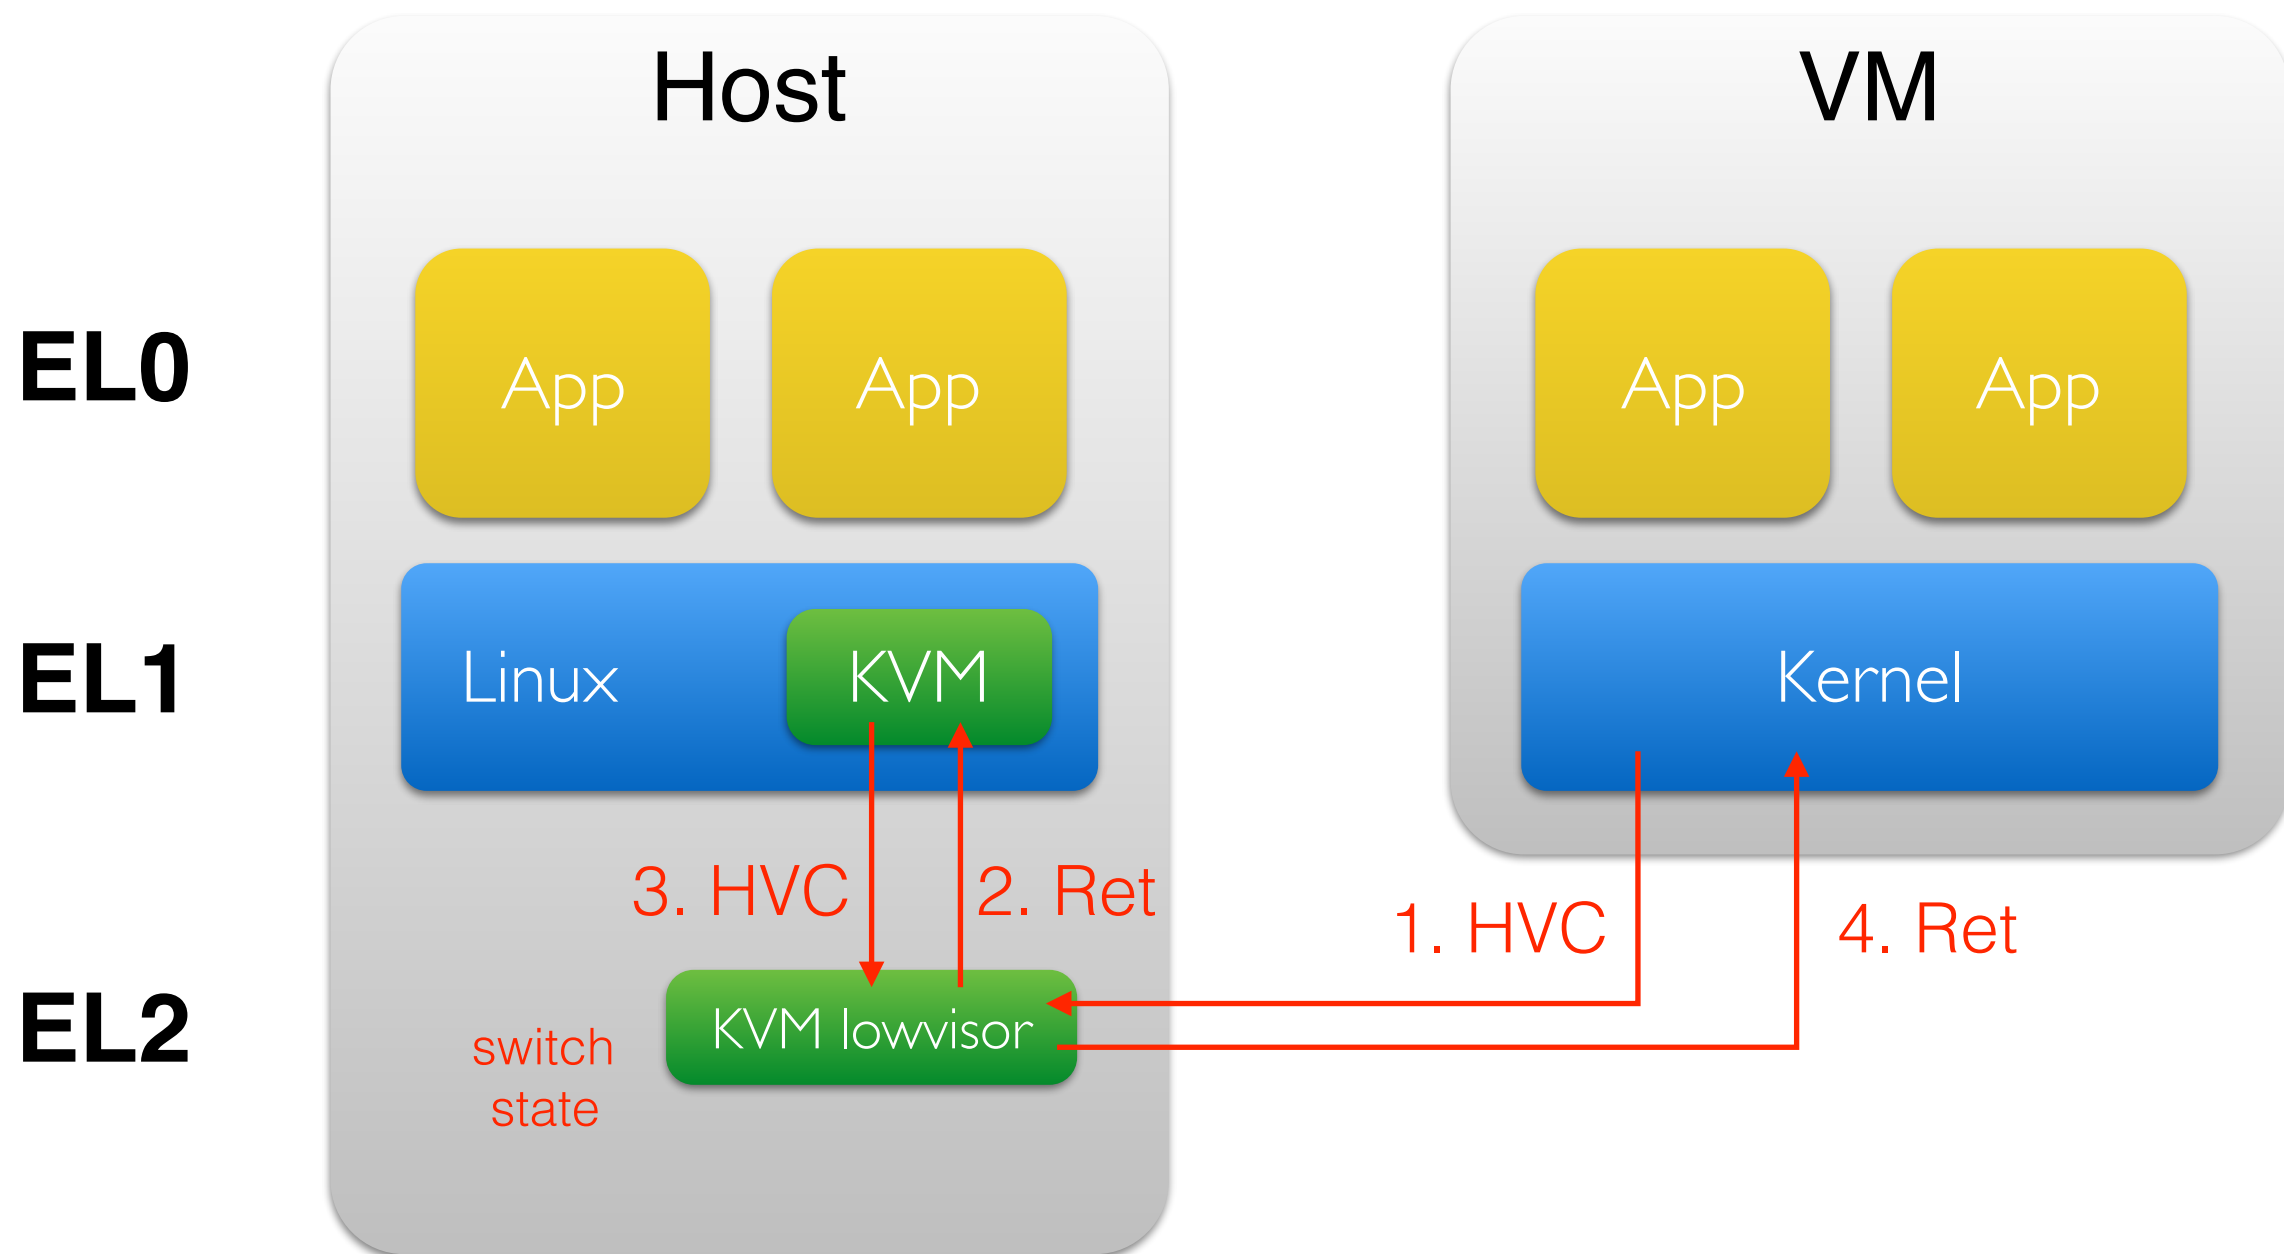
# Xen ARM: Bare-Metal



# KVM/ARM: Hosted



# KVM/ARM: Hosted



# Micro Results

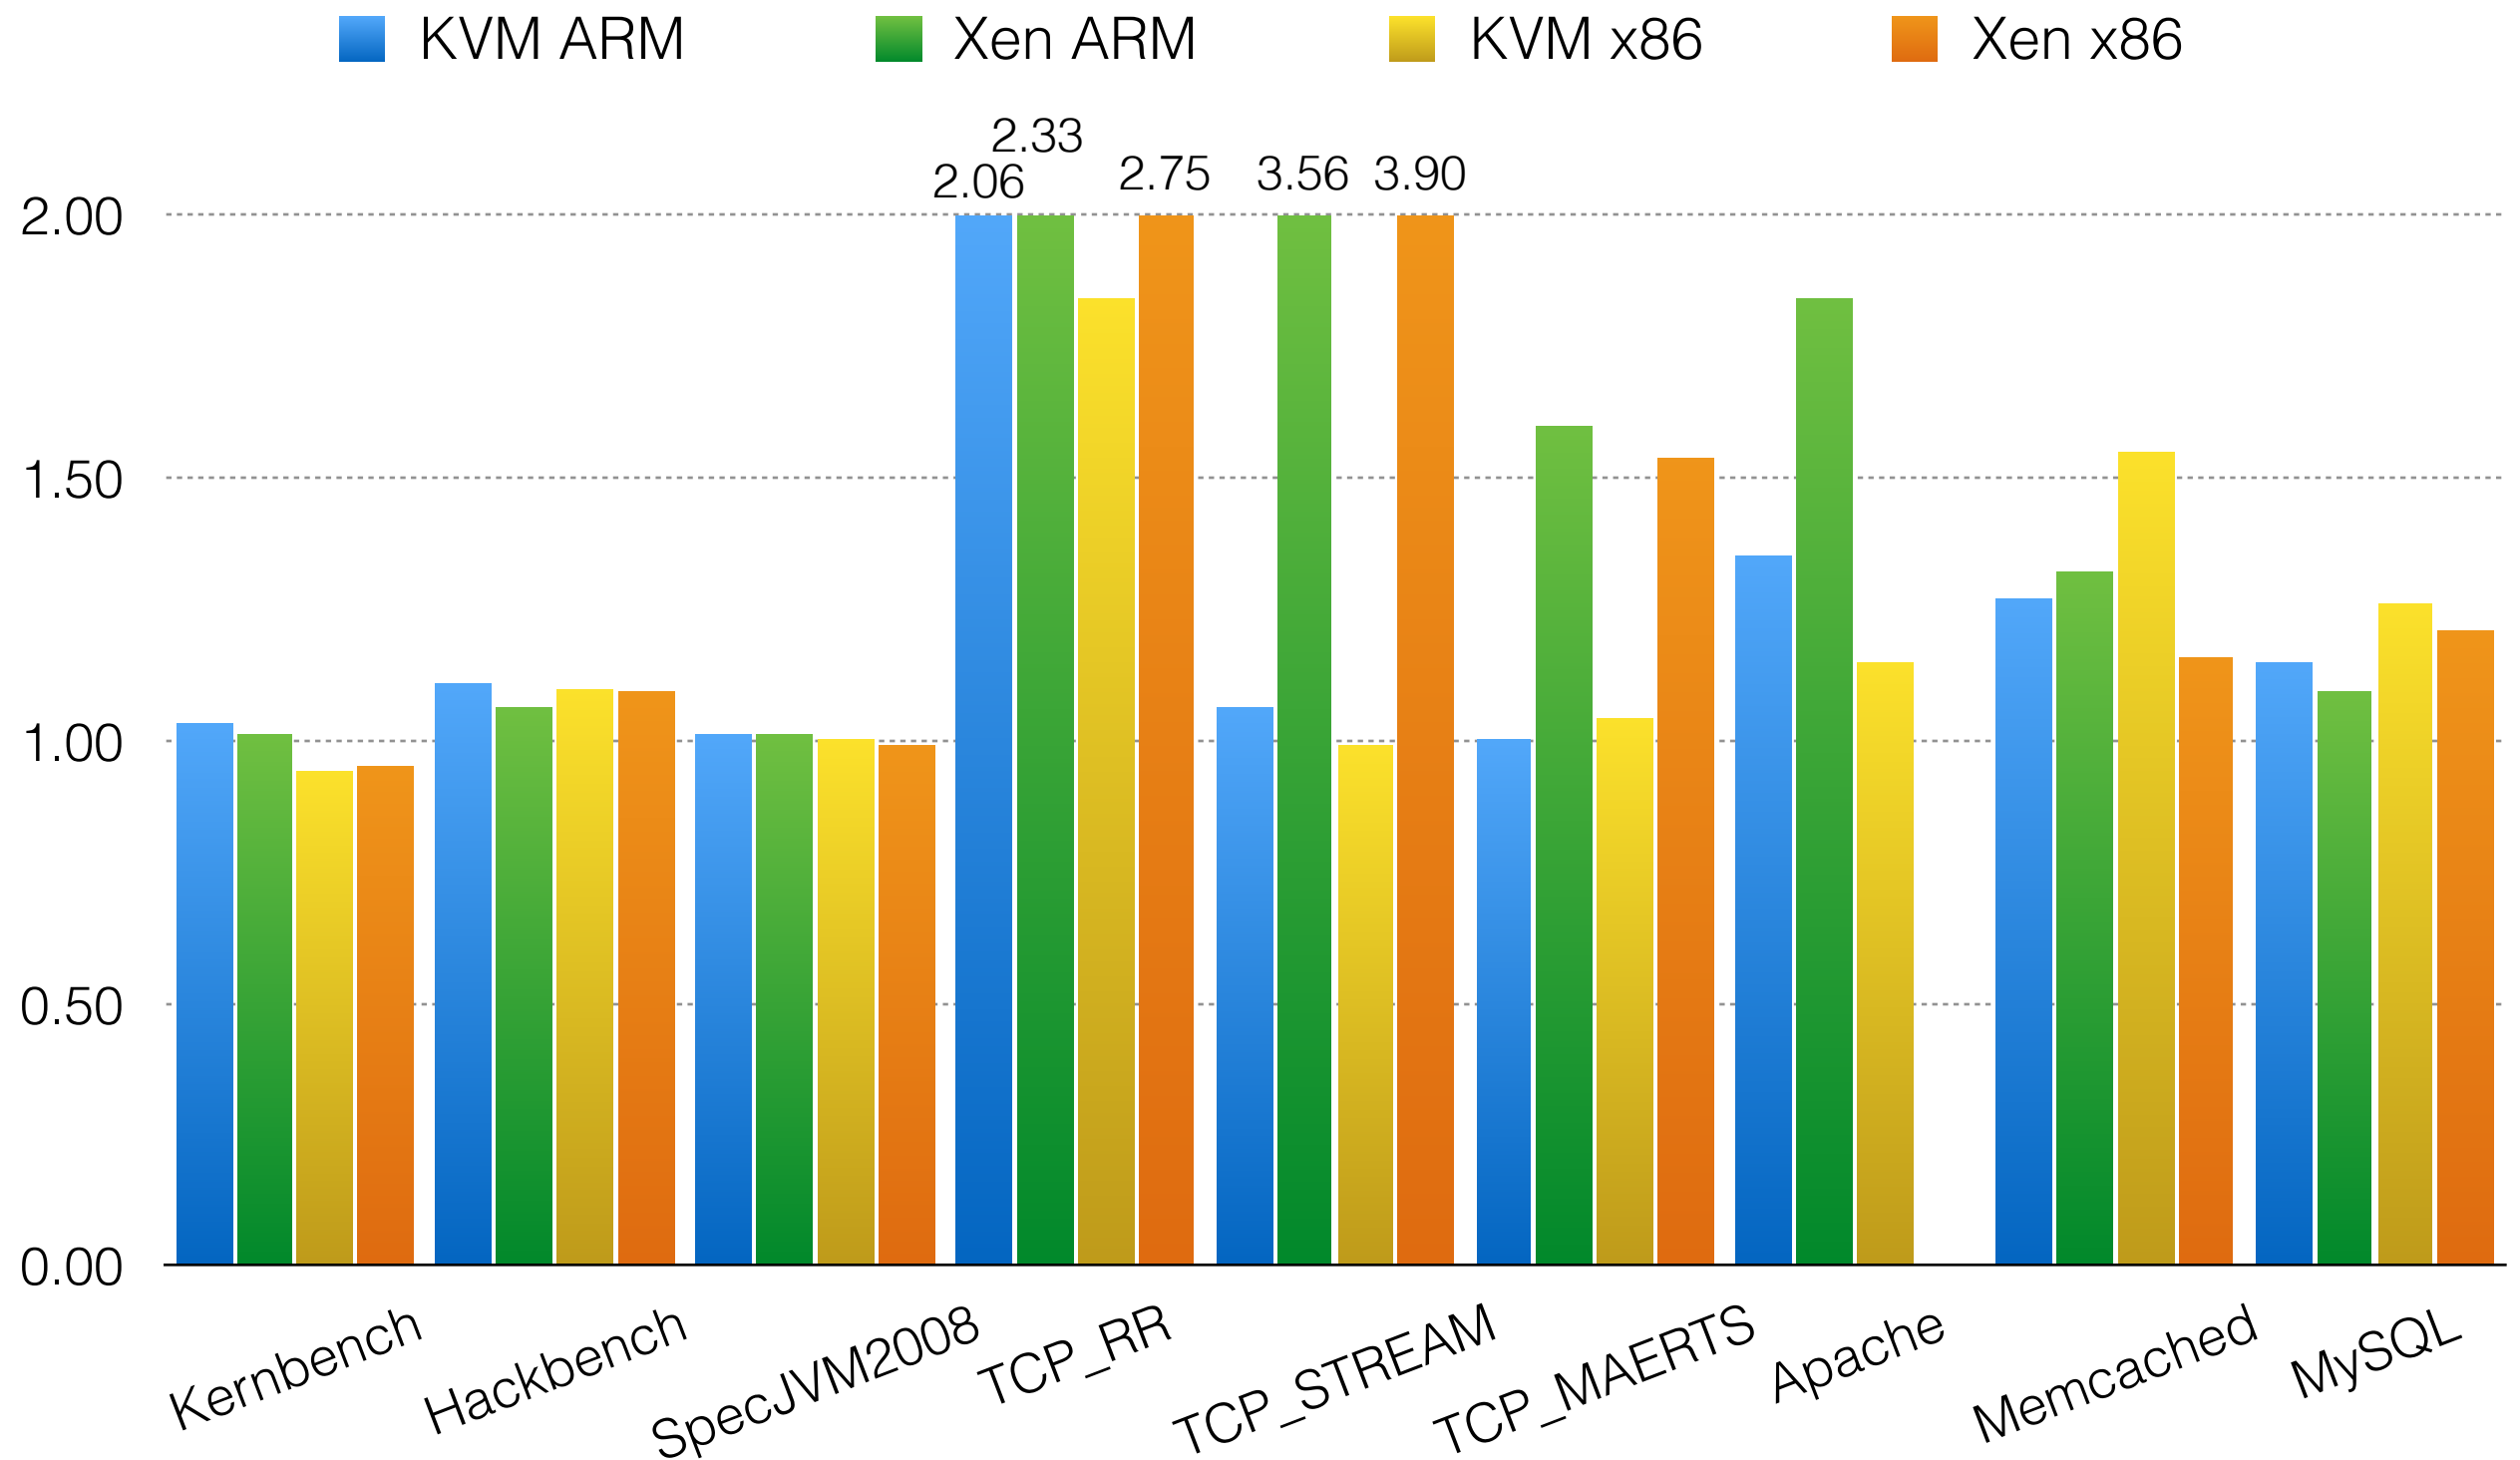
CPU Clock Cycles	ARM		x86	
	KVM	Xen	KVM	Xen
Hypercall	6,500	376	1,300	1,228

- 1: ARM can be either much faster or slower than x86
  - > x86 VM Exit more complicated than ARM Trap
- 2: KVM is much slower than Xen on ARM
  - > ARM architecture not designed for Type 2

# Application Workloads

Application	Description
Kernbench	Kernel compile
Hackbench	Scheduler stress
SPECjvm2008	Java workload
Netperf	Network performance
Apache	Web server stress
Memcached	Key-Value store
MySQL	Database workload

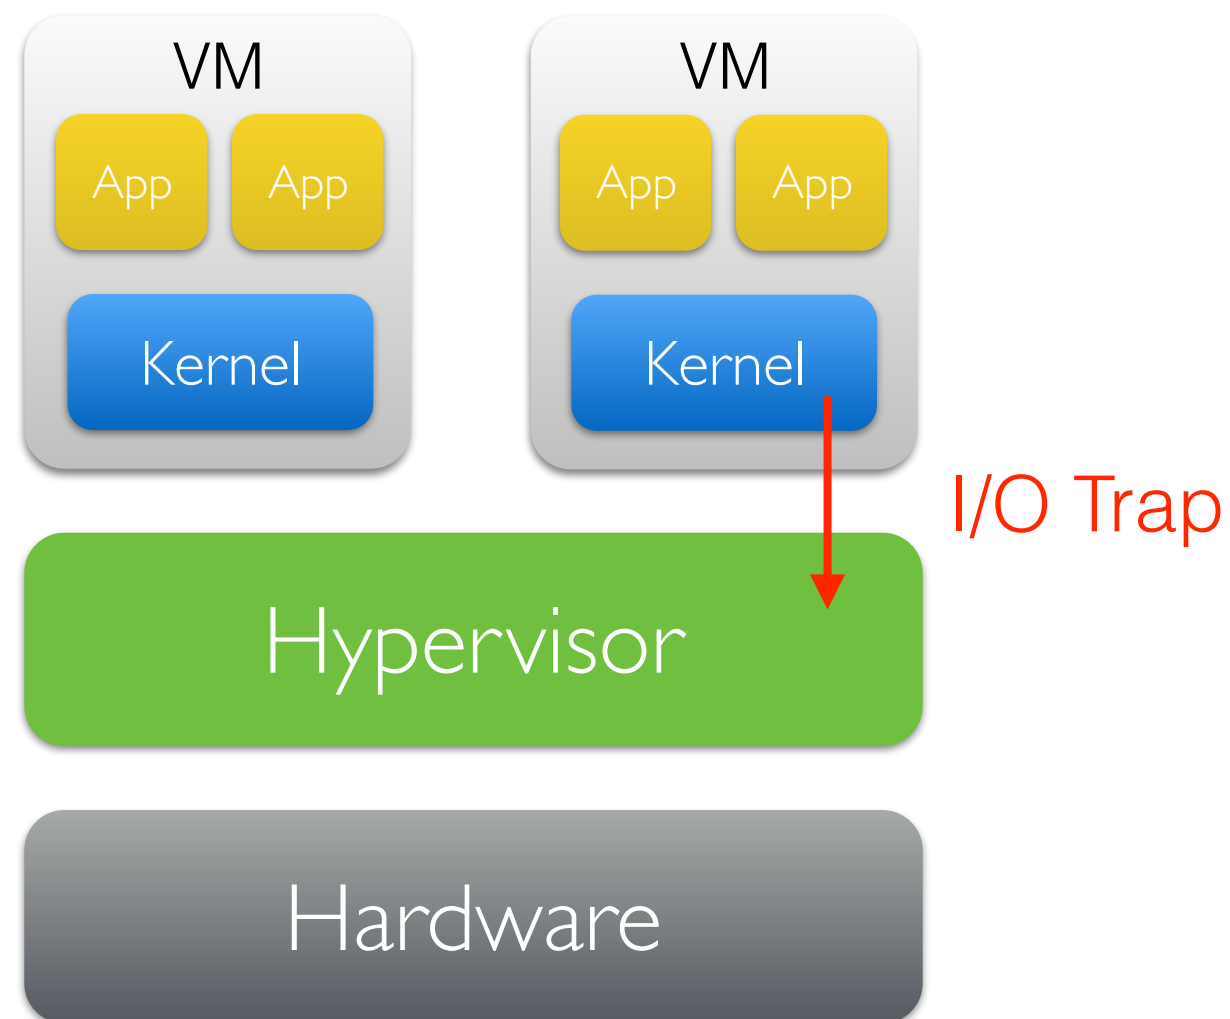
# Application Performance



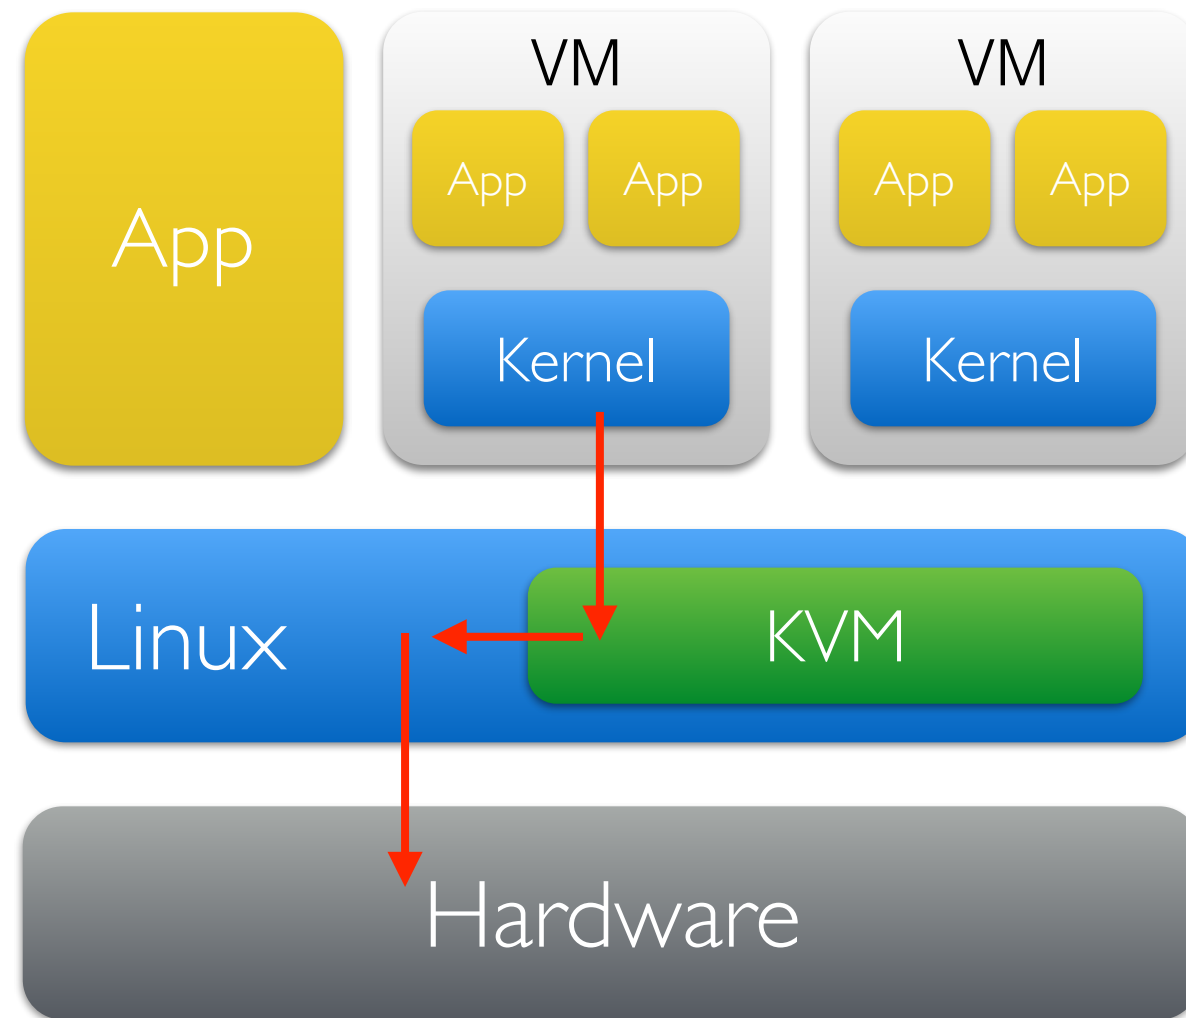
Normalized overhead (lower is better)



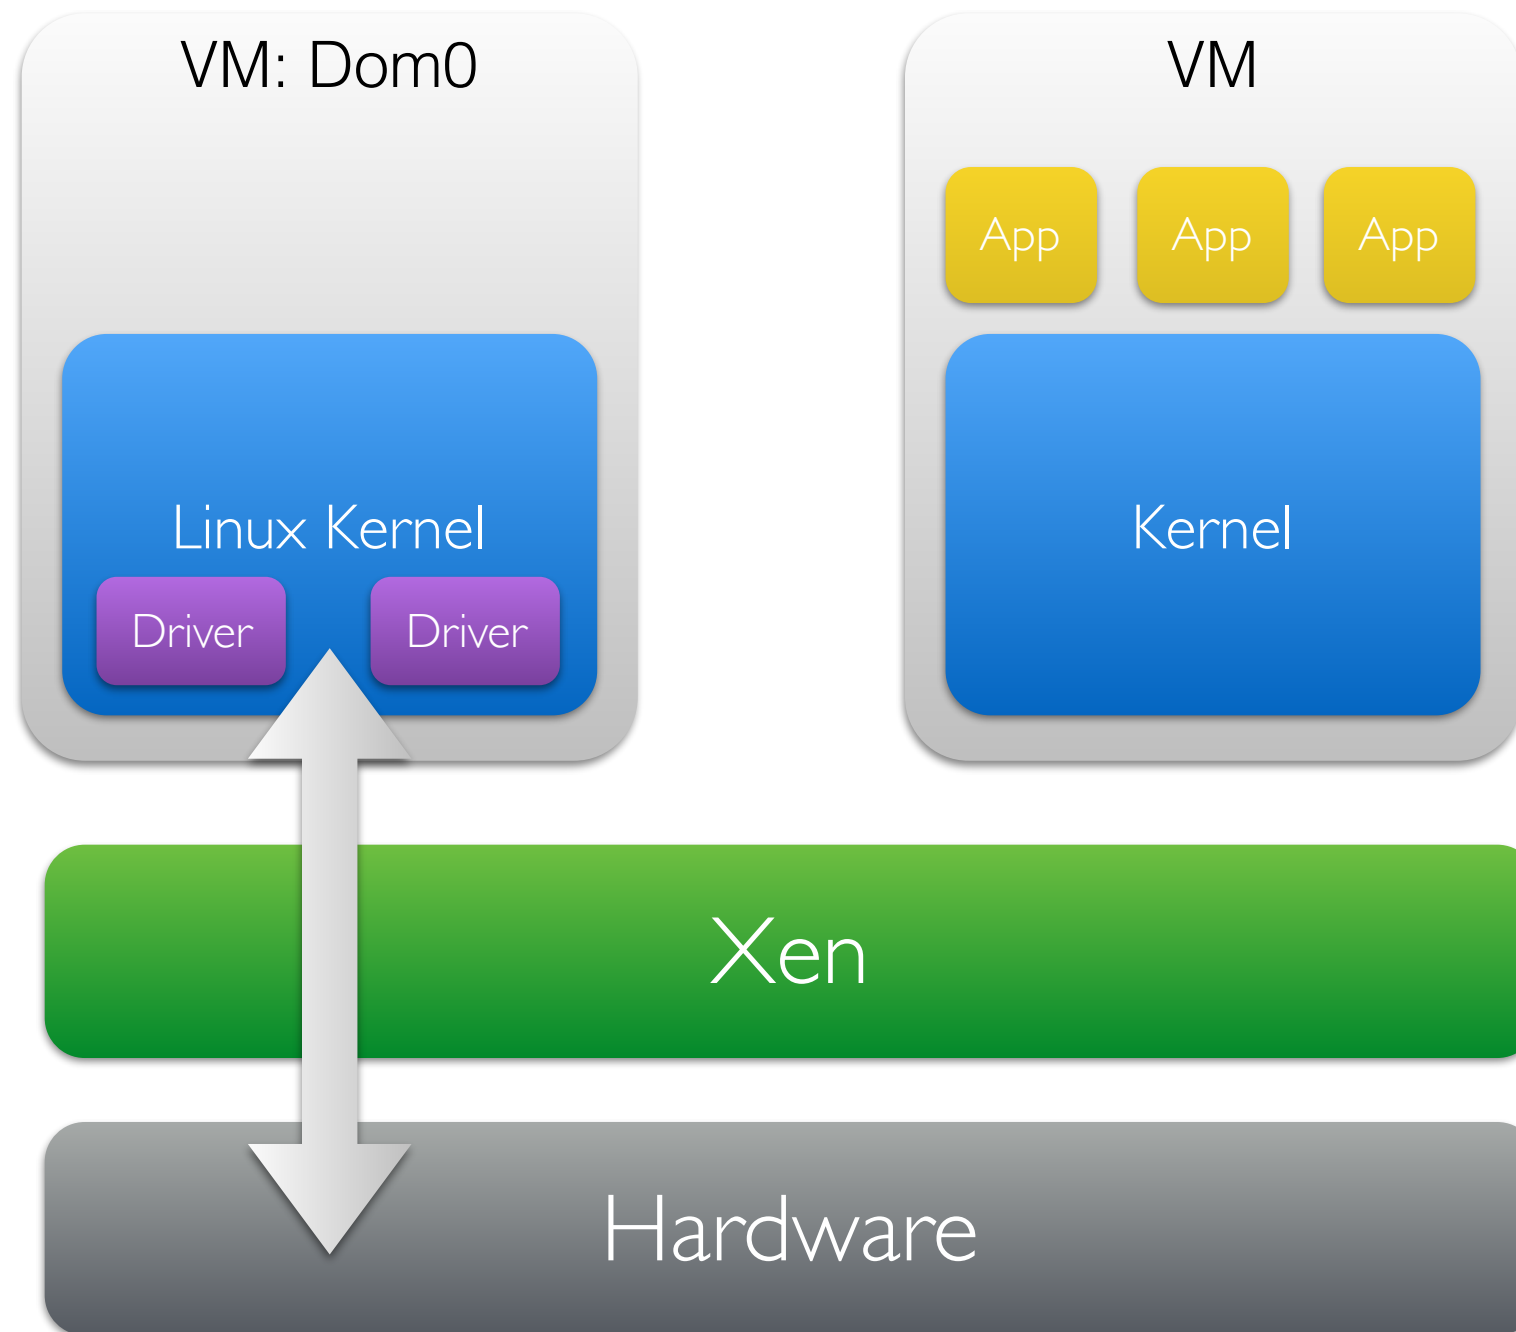
# Virtualized I/O



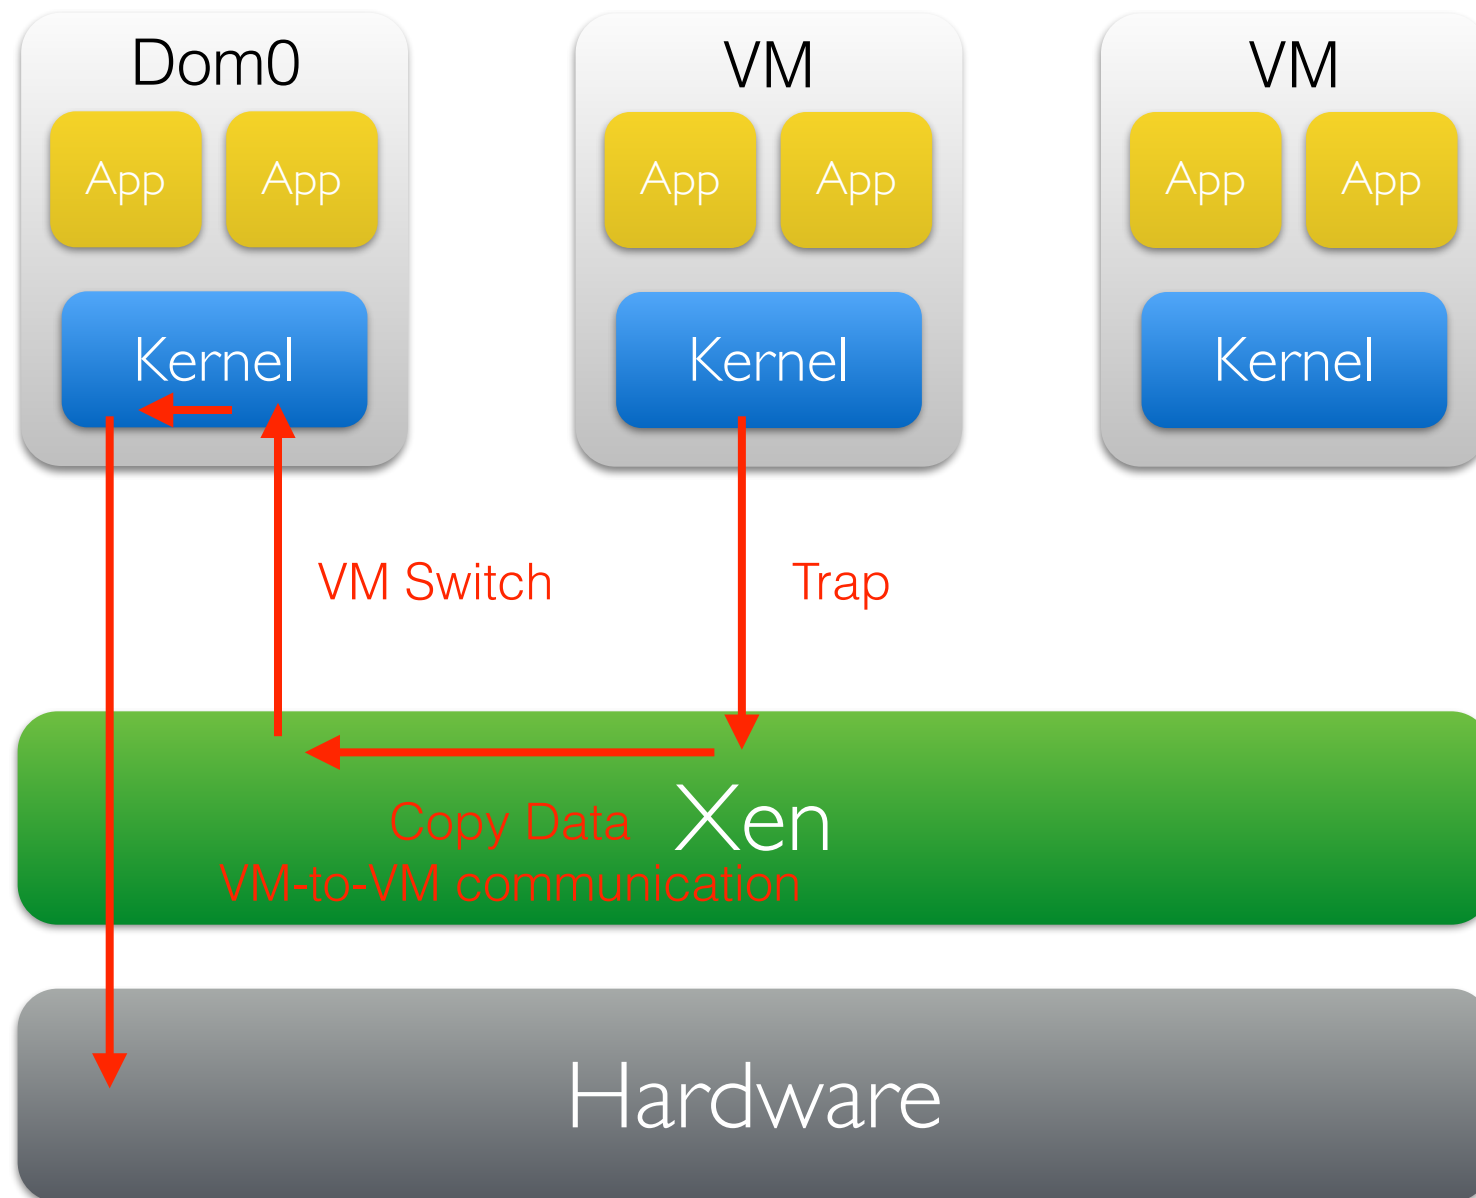
# KVM I/O Model



# Xen Device Drivers



# Xen I/O Model




# KVM ARM I/O

- Trap is slow
- But all you do is a trap

# Xen ARM I/O

- Trap is fast
- But you do much more...

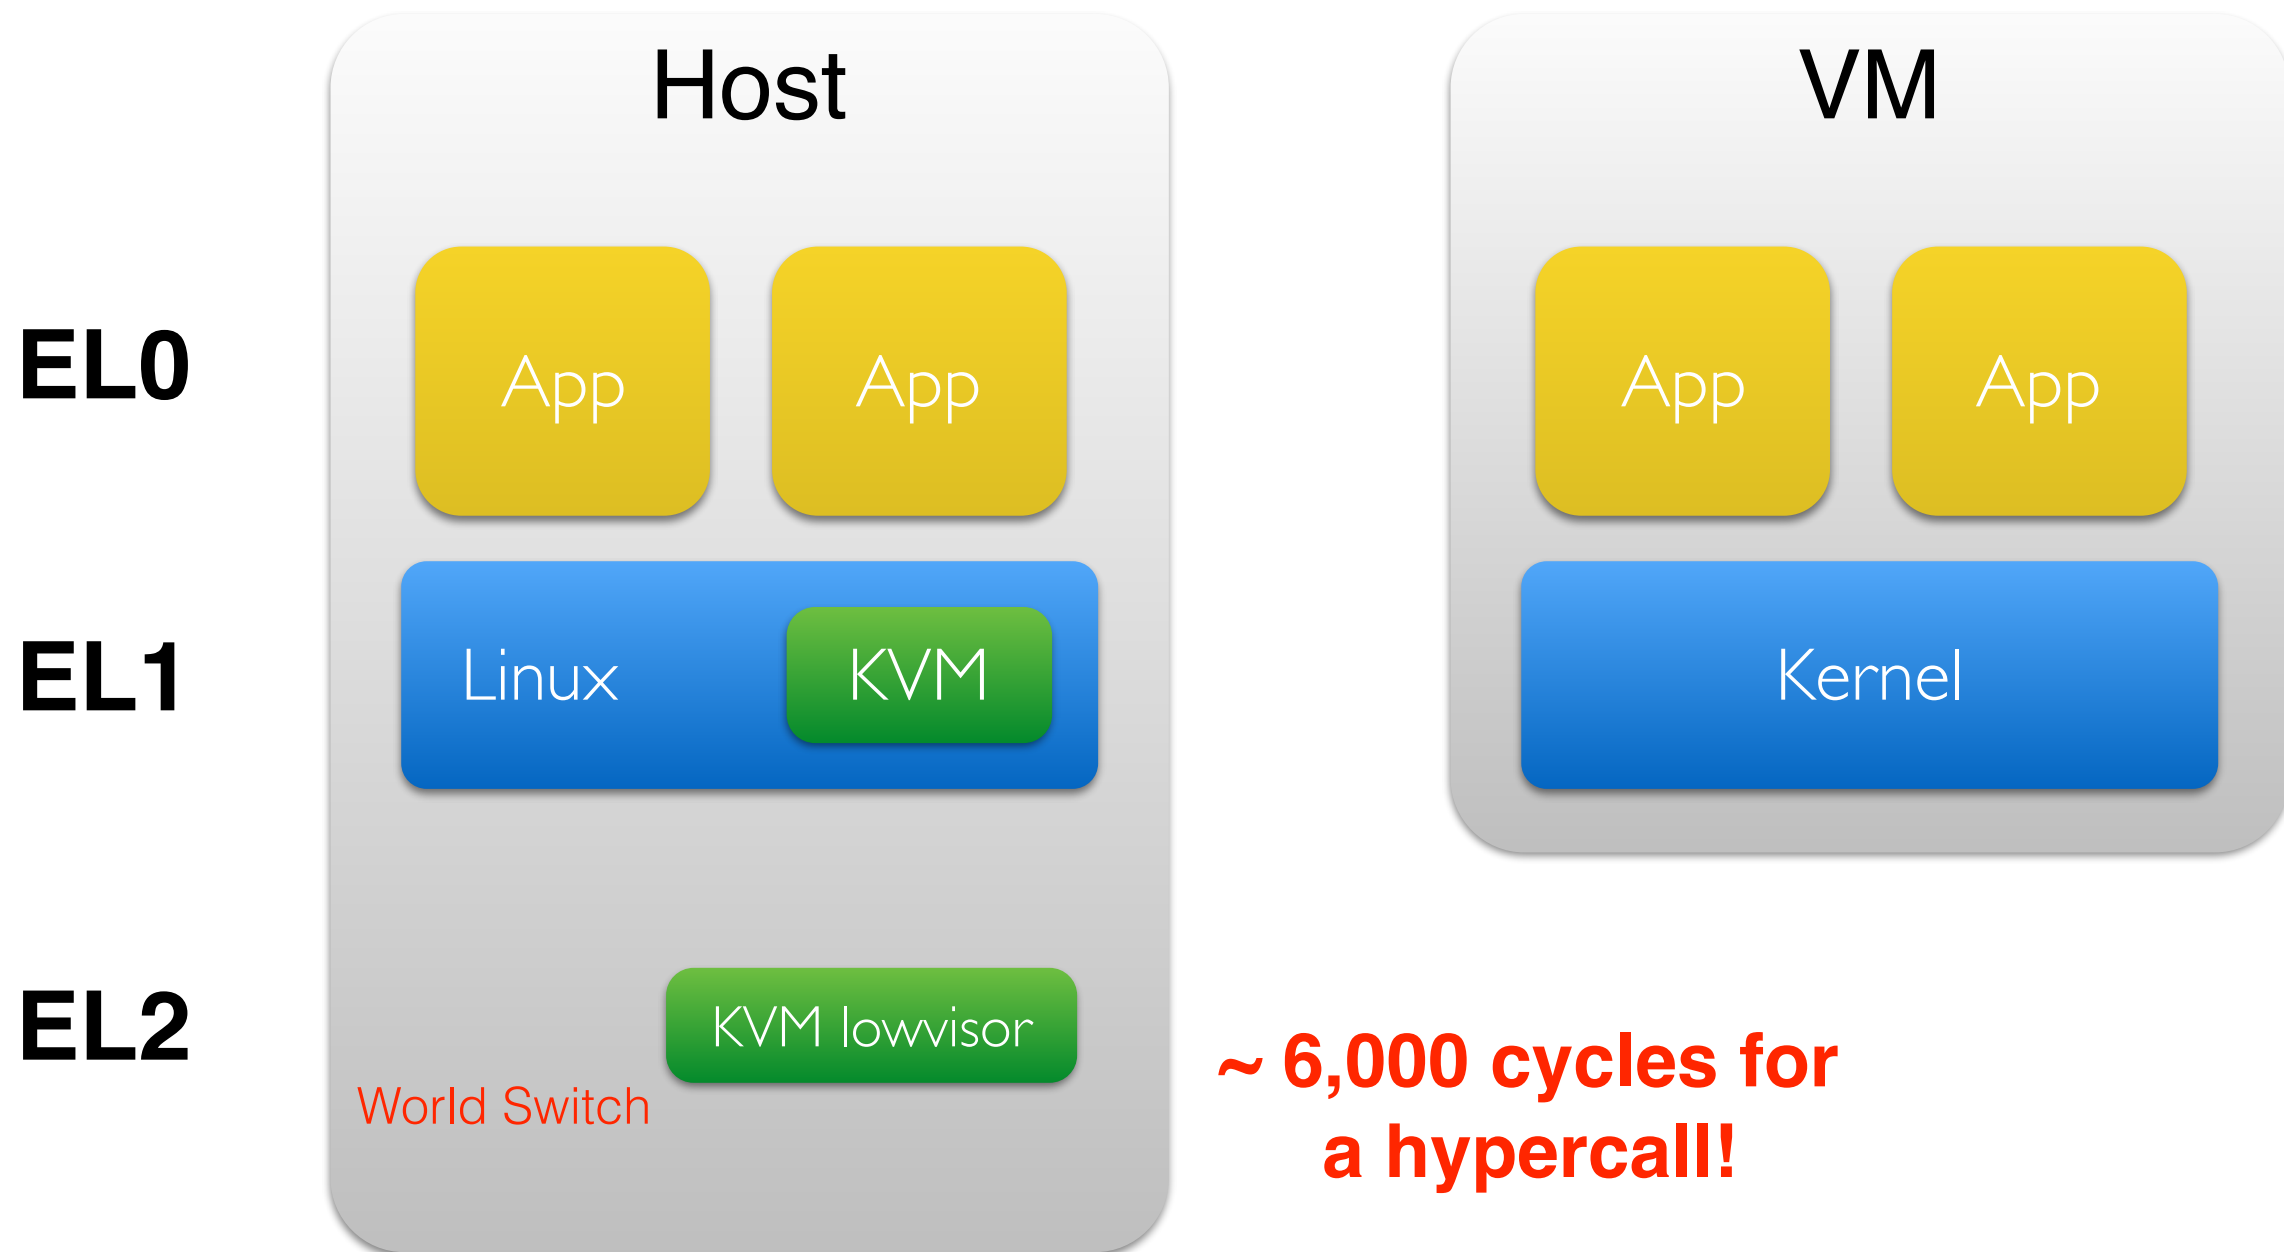
# Architecture Improvements



# VHE

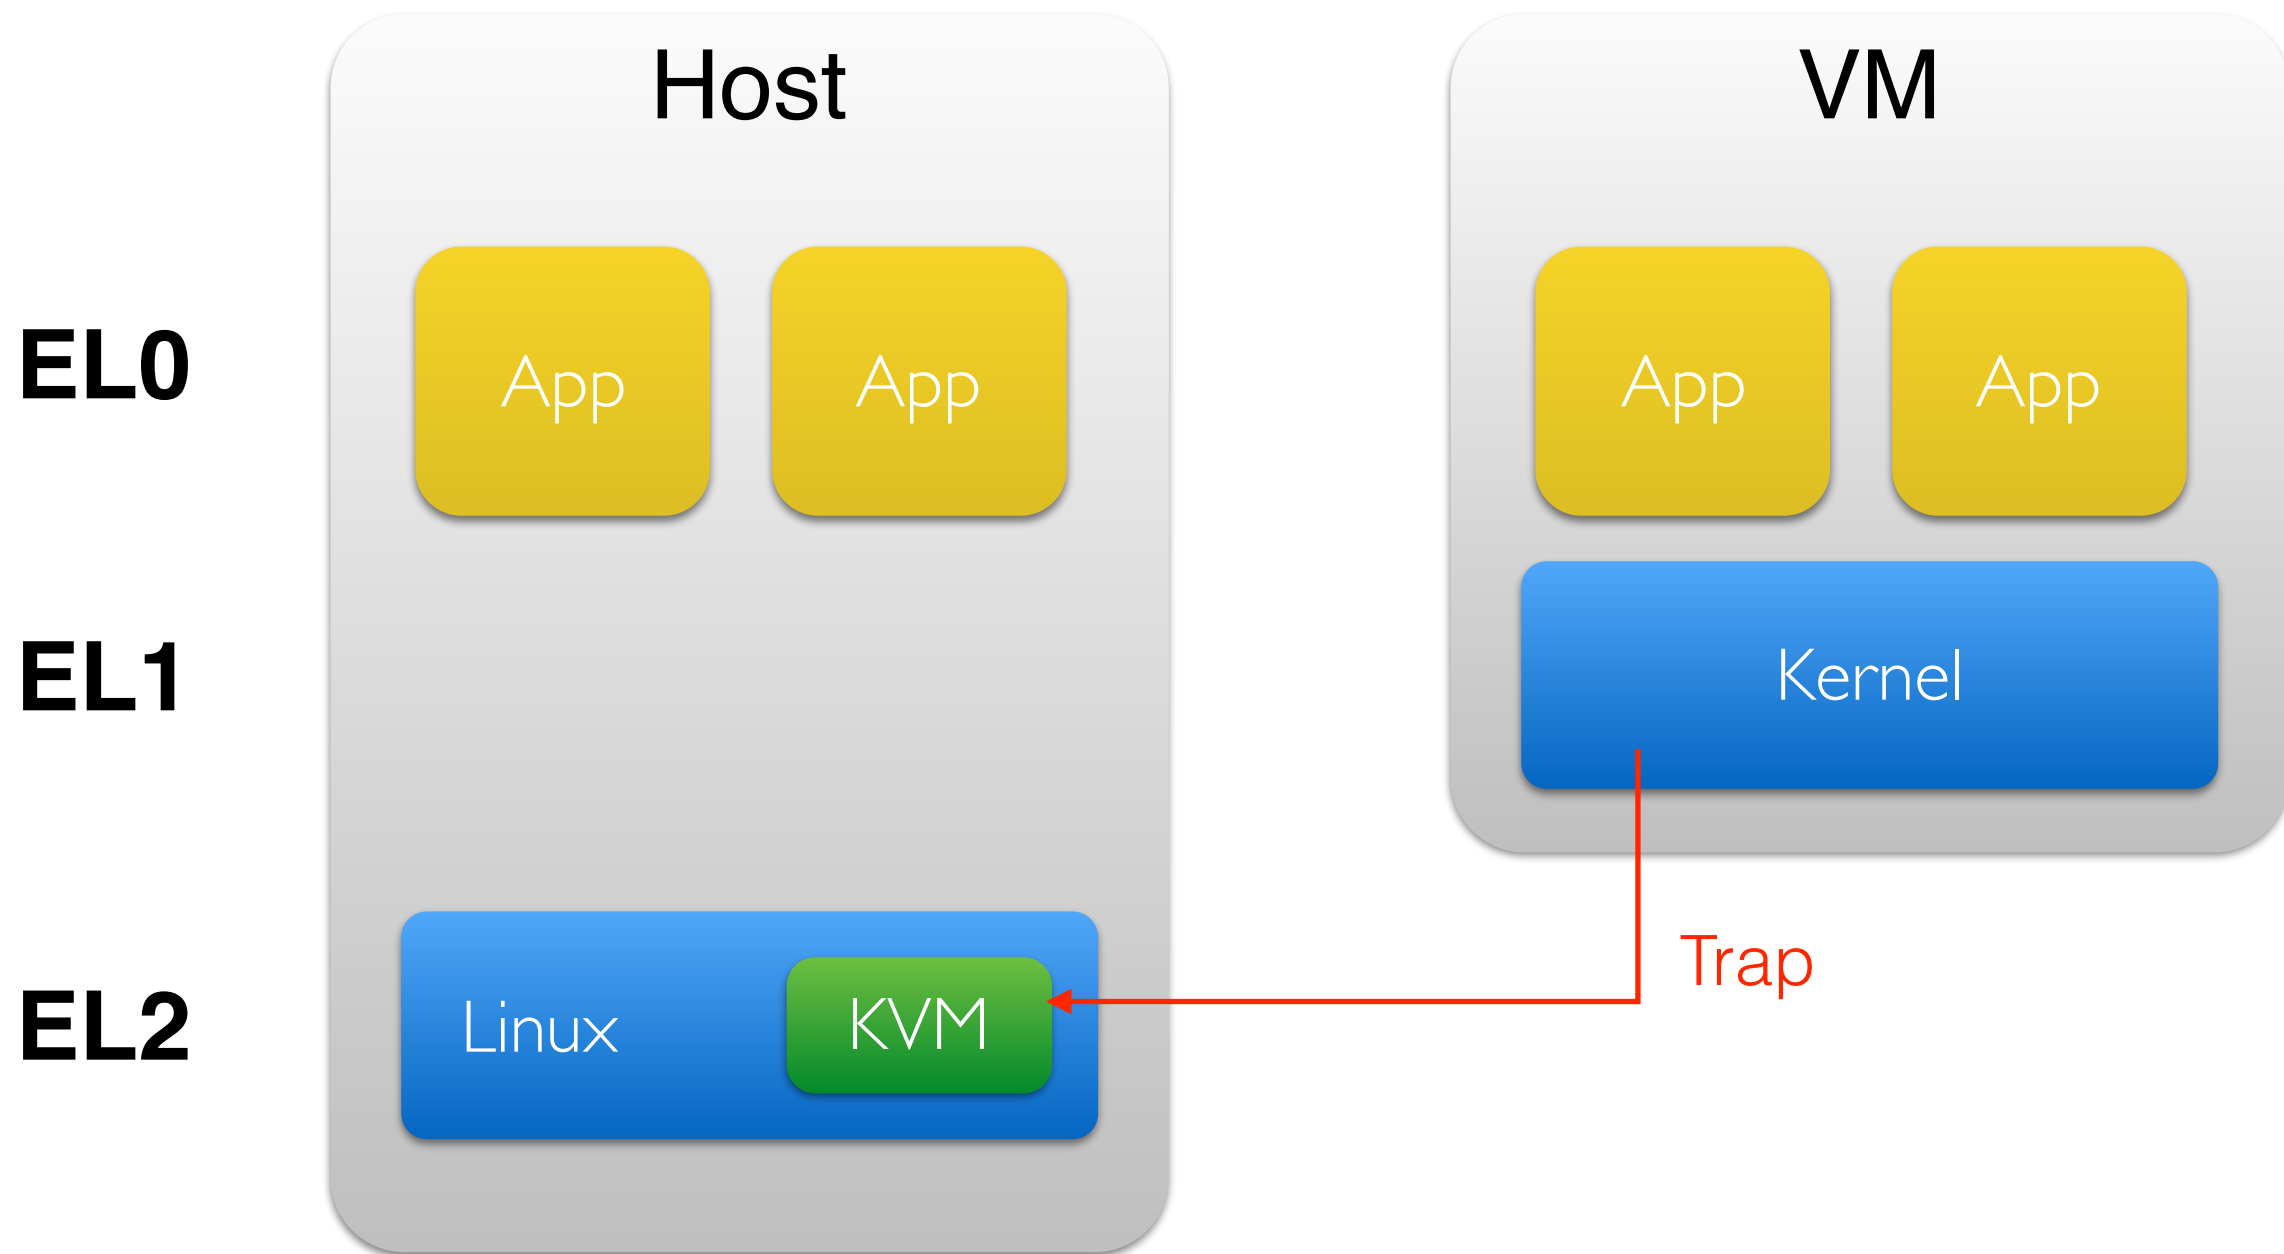
Virtualization Host Extension

# KVM





# KVM + VHE



# VHE

1. Expand EL2 to support all EL1 features
2. EL1 register accesses to go to EL2

# VHE

- Available in ARMv8.1
- No (public) hardware yet

# Conclusions

- Micro operations: ARM can be faster than x86
- Not achievable for Type 2 hypervisors
- Type 1 is dominated by other I/O costs
- ARM overhead is comparable to x86
- ARMv8.1 adds VHE for hosted hypervisors
- The software matters!