


# kraxel's news

TAG ARCHIVES: VGPU

## vgpu display support finally merged upstream

It took more than a year from the first working patches to the upstream merge. But now it's finally done. The linux kernel 4.16 (released on easter weekend) has the kernel-side code needed. The qemu code has been merged too (for gtk and spice user interfaces) and will be in the upcoming 2.12 release which is in code freeze right now. The 2.12 release candidates already have the code, so you can grab one if you don't want wait for the final release to play with this.

The vgpu code in the intel driver is off by default and must be enabled via module option. And, while being at it, also suggest to load the kvmgt module. So I've dropped a config file with these lines ...



```
options i915 enable_gvt=1
softdep i915 pre: kvmgt
```

... into `/etc/modprobe.d/`. For some reason dracut didn't pick the changes up even after regenerating the `initrd`. Because of that I've blacklisted the intel driver (`rd.driver.blacklist=i915` on the kernel command line) so the driver gets loaded later, after mounting the root filesystem, and `modprobe` actually sets the parameter.

With that in place you should have a `/sys/class/mdev_bus` directory with the intel gpu in there. You can create vgpu devices now. Check the [mediated device documentation](#) for details.

One final thing to take care of: Currently using `gvt` mlocks all guest memory. For that work the `mlock` limit (`ulimit -l`) must be high enough, otherwise the vgpu will not work correctly and you'll see a scrambled display. Limit can be configured in `/etc/security/limits.conf`.

Now lets use our new vgpu with qemu:

```
qemu-system-x86_64 \
  -enable-kvm \
  -m 1G \
  -nodefaults \
  -M graphics=off \
  -serial stdio \
  -display gtk,gl=on \
  -device vfio-pci,syfsdev=/sys/bus/mdev/devices/UUID,display=on \
  -cdrom /vmdisk/iso/Fedora-Workstation-Live-x86_64-27-1.6.iso
```

Details on the non-obvious qemu switches:

### **-nodefaults**

Do not create default devices (such as vga and nic).

### **-M graphics=off**

Hint for the firmware that the guest runs without a graphical display. This enables serial console support in seabios. We use this here because the vgpu has no firmware support (i.e. no vgabios), therefore nothing is visible on the display until the i915 kernel module loads.

### **-display gtk,gl=on**

Use gtk display, enable opengl.

### **-device vfio-pci,syfsdev=/sys/bus/mdev/devices/*UUID*,display=on**

Add the vgpu to the guest, enable the display. Of course you have to replace *UUID* with your device.

libvirt support is still being worked on. Most bits are there, but some little details are missing. For example there is no way (yet) to tell libvirt the guest doesn't need an emulated vga device, so you'll end up with two spice windows, one for the emulated vga and one for the vgpu. Other than that things are working pretty straight forward. You need spice with opengl support enabled:

```
<graphics type='spice'>
  <listen type='none' />
  <gl enable='yes' />
</graphics>
```

And the vgpu must be added of course:

```
<hostdev mode='subsystem' type='mdev' managed='no' model='vfio-pci'>
  <source>
    <address uuid='UUID' />
  </source>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x0' />
</hostdev>
```

Then you can start the domain. Use “*virt-viewer -attach guest*” to connect to the guest. Note that guest and virt-viewer must run on the same machine, sending the vgpu display to a remote machine does not yet work.

[<https://www.kraxel.org/blog/2018/04/vgpu-display-support-finally-merged-upstream/>] .


After a looong time things are finally landing upstream. Both vfio API update for vgpu local display support and support for that in the intel drm driver have been merged during the 4.16 merge window. Should not take too long to get the qemu changes merged upstream now, if all goes well the next release (2.12) will have it.

This entry was posted in Tech Tips and tagged qemu, vgpu on February 13, 2018

[<https://www.kraxel.org/blog/2018/02/local-vgpu-display-status-update/>].

[illegible]

<https://www.kraxel.org/blog/tag/vgpu/>



host kernel: <https://www.kraxel.org/cgit/linux/log/?h=intel-vgpu>  
qemu: <https://www.kraxel.org/cgit/qemu/log/?h=work/intel-vgpu>

This entry was posted in Uncategorized and tagged qemu, vgpu on January 24, 2017 [https://www.kraxel.org/blog/2017/01/local-display-for-intel-vgpu-starts-working/].

The upstreaming process of virtual gpu support (vgpu) made a big step forward with the 4.10 merge window. Two important pieces have been merged:

Second, intel landed initial mdev support for the i915 driver too. There is quite some work left to do in future kernel releases though. Accessing to the guest display is not supported yet, so you must run `x11vnc` or simliar tools in the guest to see the screen. Also there are some stability issues to find and fix.

## On the host: create virtual devices

<https://www.kraxel.org/blog/tag/vgpu/>

CONFIG\_DRM\_I915\_GVT and CONFIG\_DRM\_I915\_GVT\_KVMGT for intel vgpu support. Building the mttty sample driver (CONFIG\_SAMPLE\_VFIO\_MDEV\_MTTY, a virtual serial port) can be useful too, for testing.

Boot the new kernel. Load all modules: vfio-pci, vfio-mdev, optionally mttty. Also i915 and kvmgt of course, but that probably happened during boot already.

Go to the /sys/class/mdev\_bus directory. This should look like this:

```
kraxel@broadwell ~# cd /sys/class/mdev_bus
kraxel@broadwell .../class/mdev_bus# ls -l
total 0
lrwxrwxrwx. 1 root root 0 17. Jan 10:51 0000:00:02.0 -> ../../devices/pci0000:00/0000:00:
lrwxrwxrwx. 1 root root 0 17. Jan 11:57 mttty -> ../../devices/virtual/mtty/mtty
```

Each driver with mdev support has a directory there. Go to `$device/mdev_supported_types` to check what kind of virtual devices you can create.

```
kraxel@broadwell .../class/mdev_bus# cd 0000:00:02.0/mdev_supported_types
kraxel@broadwell .../0000:00:02.0/mdev_supported_types# ls -l
total 0
drwxr-xr-x. 3 root root 0 17. Jan 11:59 i915-GVTg_V4_1
drwxr-xr-x. 3 root root 0 17. Jan 11:57 i915-GVTg_V4_2
drwxr-xr-x. 3 root root 0 17. Jan 11:59 i915-GVTg_V4_4
```

As you can see intel supports three different configurations on my machine. The configuration (basically the amount of video memory) differs, and the number of instances you can create. Check the `description` and `available_instance` files in the directories:

```
kraxel@broadwell .../0000:00:02.0/mdev_supported_types# cd i915-GVTg_V4_2
kraxel@broadwell .../mdev_supported_types/i915-GVTg_V4_2# cat description
low_gm_size: 64MB
high_gm_size: 192MB
fence: 4
kraxel@broadwell .../mdev_supported_types/i915-GVTg_V4_2# cat available_instance
2
```

Now it is possible to create virtual devices by writing a UUID into the `create` file:

```
kraxel@broadwell .../mdev_supported_types/i915-GVTg_V4_2# uuid=$(uuidgen)
kraxel@broadwell .../mdev_supported_types/i915-GVTg_V4_2# echo $uuid
f321853c-c584-4a6b-b99a-3eee22a3919c
kraxel@broadwell .../mdev_supported_types/i915-GVTg_V4_2# sudo sh -c "echo $uuid > create"
```

The new vgpu device will show up as subdirectory of the host gpu:

```
kraxel@broadwell .../mdev_supported_types/i915-GVTg_V4_2# cd ../../$uuid
kraxel@broadwell .../0000:00:02.0/f321853c-c584-4a6b-b99a-3eee22a3919c# ls -l
total 0
lrwxrwxrwx. 1 root root    0 17. Jan 12:31 driver -> ../../../../../../bus/mdev/drivers/vfio_md
lrwxrwxrwx. 1 root root    0 17. Jan 12:35 iommu_group -> ../../../../../../kernel/iommu_groups
lrwxrwxrwx. 1 root root    0 17. Jan 12:35 mdev_type -> ../mdev_supported_types/i915-GVTg
drwxr-xr-x. 2 root root    0 17. Jan 12:35 power
--w-----, 1 root root 4096 17. Jan 12:35 remove
lrwxrwxrwx. 1 root root    0 17. Jan 12:31 subsystem -> ../../../../../../bus/mdev
-rw-r--r--, 1 root root 4096 17. Jan 12:35 uevent
```

You can see the device landed in iommu group 10. We'll need that in a moment.

### On the host: configure guests

Ideally this would be as simple as adding `<hostdev>` to your guests libvirt xml config. The mdev devices don't have a pci address on the host though, and because of that they must be passed to qemu using the sysfs device path instead of the pci address. libvirt doesn't (yet) support sysfs paths though, so it is a bit more complicated for now. A lot of the setup libvirt does for hostdevs automatically must be done manually instead.

First, we must allow qemu access /dev. By default libvirt uses control groups to restrict access. That must be turned off. Edit `/etc/libvirt/qemu.conf`. Uncomment the `cgroup_controllers` line. Remove "devices" from the list. Restart libvirtd.

Second, we must allow qemu access the iommu group (10 in my case). A simple `chmod` will do:

```
kraxel@broadwell ~# chmod 666 /dev/vfio/10
```

Third, we must update the guest configuration:

```
<domain type='kvm' xmlns:qemu='http://libvirt.org/schemas/domain/qemu/1.0'>
[ ... ]
```

```

<currentMemory unit='KiB'>1048576</currentMemory>
<memoryBacking>
  <locked/>
</memoryBacking>
[ ... ]
<qemu:commandline>
  <qemu:arg value='-device' />
  <qemu:arg value='vfio-pci,addr=05.0,sysfsdev=/sys/class/mdev_bus/0000:00:02.0/f3' />
</qemu:commandline>
</domain>

```

There is special qemu namespace which can be used to pass extra command line arguments to qemu. We do this here to use a qemu feature not yet supported by libvirt (use sysfs paths for vfio-pci). Also we must explicitly allow to lock down guest memory.

Now we are ready to go:

```
kraxel@broadwell ~# virsh start --console $guest
```

## In the guest

It is a good idea to prepare the guest a bit before adding the vgpu to the guest configuration. Setup a serial console, so you can talk to it even in case graphics are broken. Blacklist the i915 module and load it manually, at least until you have a known-working configuration. Also booting to runlevel 3 (aka multi-user.target) instead of 5 (aka graphical.target) and starting the xorg server manually is better for now.

For the guest machine intel recommends the 4.8 kernel. In theory newer kernels should work too, in practice they didn't last time I tested (4.10-rc2). Also make sure the xorg server uses the modesetting driver, the intel driver didn't work in my testing. This config file will do:

```

root@guest ~# cat /etc/X11/xorg.conf.d/intel.conf
Section "Device"
    Identifier "Card0"
#    Driver "intel"
    Driver "modesetting"
    BusID "PCI:0:5:0"
EndSection

```

I'm starting the xorg server with x11vnc, xterm and mwm (motif window manager) using this little script:

```

#!/bin/sh

# debug
echo "# $0: DISPLAY=$DISPLAY"

# start server
if test "$DISPLAY" != ":4"; then
    echo "# $0: starting Xorg server"
    exec startx $0 -- /usr/bin/Xorg :4

```

```
        exit 1
fi
echo "# $0: starting session"

# configure session
xrdb $HOME/.Xdefaults

# start clients
x11vnc -rfbport 5904 &
xterm &
exec mwm
```

The session runs on display 4, so you should be able to connect from the host this way:

```
kraxel@broadwell ~# vncviewer $guest_ip:4
```

Have fun!

This entry was posted in Tech Tips and tagged qemu, vgpu on January 17, 2017

[<https://www.kraxel.org/blog/2017/01/virtual-gpu-support-landing-upstream/>] .

---