

Reconnaissance of Virtio: What's new and how it's all connected?

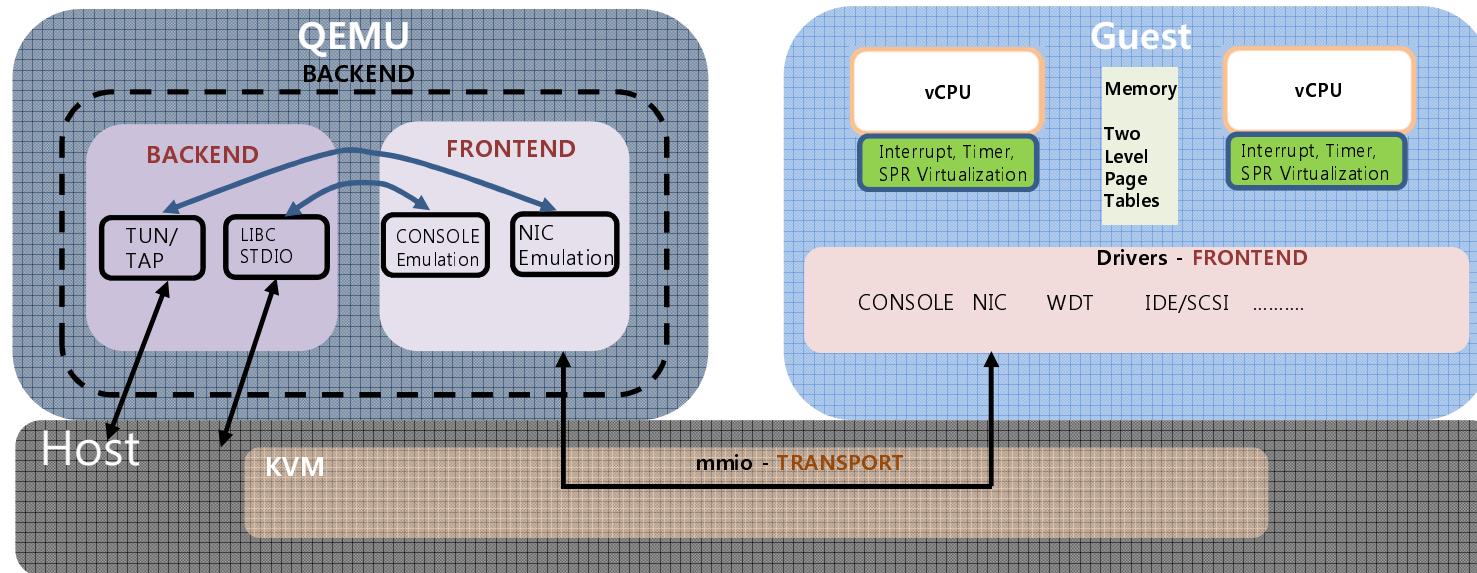
Mario Smarduch
Senior Virtualization Architect
Open Source Group
Samsung Research America (Silicon Valley)
m.smarduch@samsung.com



Agenda

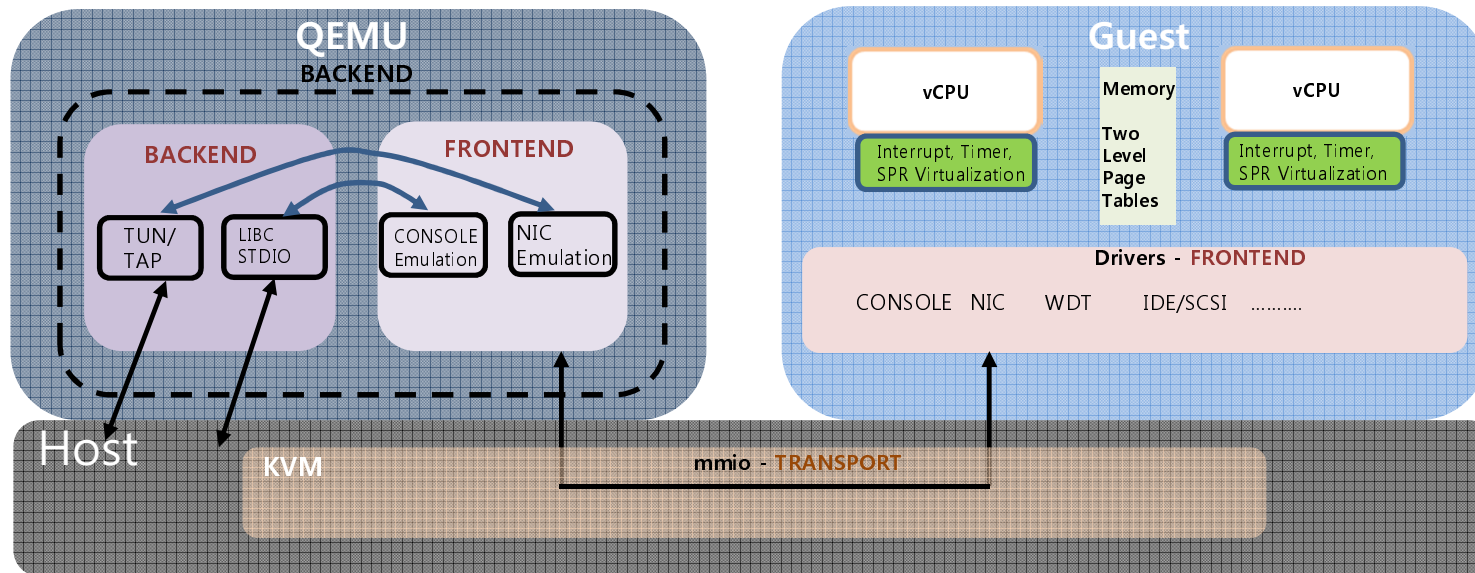
- **QEMU/Guest Machine Model & IO Overview**
- **Concepts - transport/backend – recent re-factoring**
- **PCI transport and most recent virtio-mmio transport**
- **Virtio and Device Passthrough, virtio performance**

Machine Model



- Like host, unmodified guest expects real hardware
- Machine model – combination of hw extensions, KVM, QEMU, GUEST
 - **Interrupt Local and Distributor** – hw virt extensions + kvm
 - **Special Purpose Register** – i.e. enable/disable MMU, discover CPU features – hw virt ext + kvm
 - **Timer** – hw virt extensions + kvm
 - **Memory** – hw virt + kvm
 - **Drivers/Devices** – (i) mmio (ii) para-virtualized (iii) dev passthrough
 - **Machine Model** - defines hw – CPU, Peripherals, HW address map
- Some Terms
 - **Transport** – way Guest to (i) probe, discover backend – resources; (ii) configure backend
 - **Frontend** – guest driver
 - **Backend** – whole QEMU I/O emulation + host device

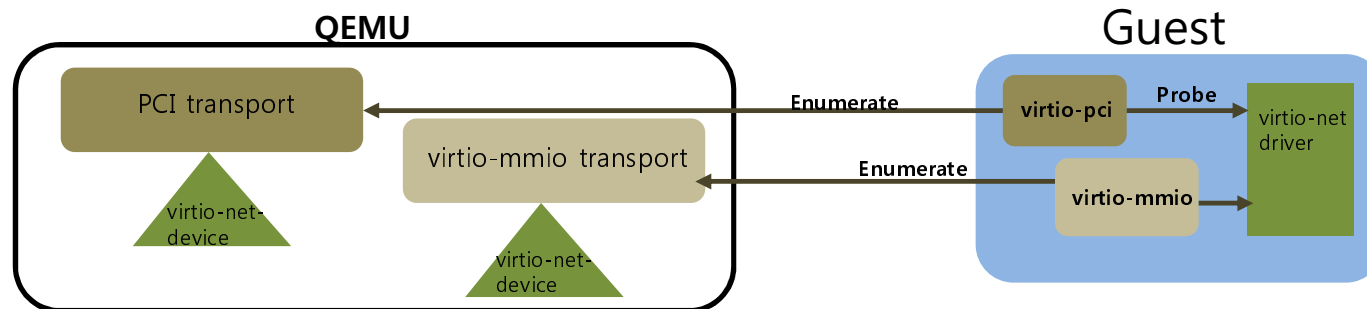
MMIO Example



- Typing a character – '-nographic'
 - Keyboard stroke – QEMU backend (IO thread) reads from stdio
 - Finds Qemu Frontend – console emulation device passes character
 - Console device injects interrupt via KVM, guest exit/resume
 - Console interrupt handler – mmio read of device buffer
 - guest exits, decodes regs to packages addr/data size
 - Returns from vCPU KVM_RUN loop to QEMU
 - QEMU finds console device handler from addr (GPA)
 - Console handler returns data at address
 - Return to KVM, data placed in dest register
 - Resume guest
- MMIO a lot of overhead!


Vision and Practice

- QEMU/Guest - Vision
 - ❑ Portability any backend plugs into any transport – no clue about transport
 - Typically one transport configured
 - '- virtio_XXX_device' option – no hint of transport – plug into first available one
 - ❑ Guest virtio driver unaware of transport
 - All transports can probe, discover backend
 - Indirect transport interface – i.e. virtio-net does not know what transport
 - ❑ Example




Vision and Practice

- In Practice – PCI preferred transport
 - ❑ Transport/backends 'fused'
 - ❑ Backend plugged into PCI
 - ❑ Prior knowledge of machine model required
 - Command line – specify transport
 - No Portability



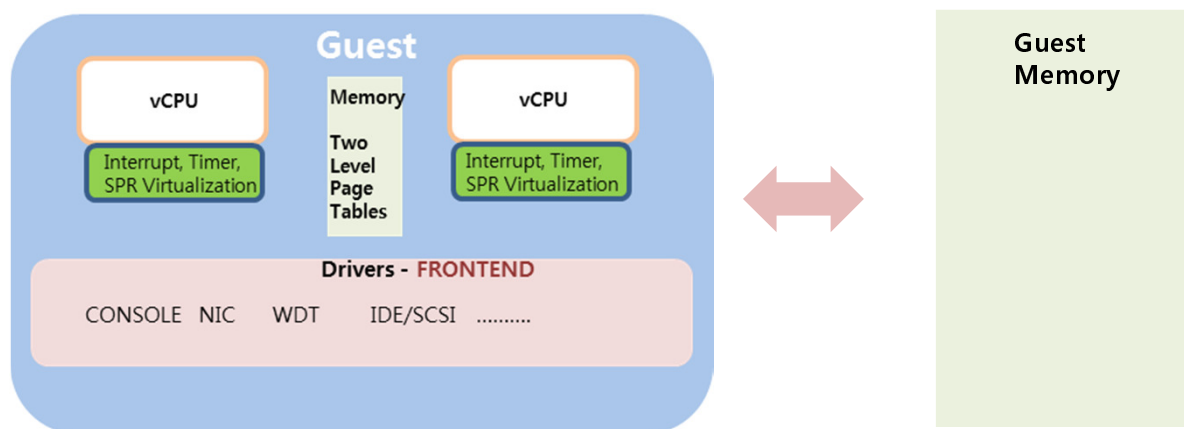
Virtio – moving data

- virtio – ring buffers accessed from several contexts
- Must deal with different addresses when moving data to/from virtio device




- Between Guest & QEMU – QEMU view
 - Host mmap() address – QEMU VA – HVA
 - To get HVA from GPA
 - Find memory region section
 - Offset = GPA – MemoryRegion base
 - Add HVA base in RAMBlock add offset
 - To get GPA from HVA
 - From RAMBlock find MemoryRegion
 - Offset = HVA address – HVA base
 - Add to MemoryRegion base address

Virtio – moving data



- Between Guest & QEMU or host – Guest view
 - Guest knows nothing about HVA
 - Current hw supports two level page tables
 - 2nd level page table maps GPA → HPA

Virtio – moving data



- Performance achieved through direct memory access (see Rusty Russels spec)

GUEST

Simple pkt – no fragments

- 1 - virtio_net_hdr (skb->cb[])
- 2 - skb->data[]


Scatterlist

- 1 – page_link = page – of virtio_net_hdr
offset = offset within page
length = sizeof virtio_net_hdr
- 2 – page_link = page – of skb->data
offset = ...
length = skb->len

Vring descriptor

- 1 – GPA addr of virtio_net_Hdr
length
 - 2 – GPA addr of skb->data
length
- NOTIFY


Virtio – moving data



QEMU

VirtQueueElement

1. `out_addr = GPA virtio_net_hdr`
`out_sg.iov_len = virtio_net_hdr length`
2. `out_addr = GPA skb->data`
`out_sg.iov_len = skb->len`



VirtQueueElement

1. `out_sg.iov_base = HVA virtio_net_hdr`
`out_sg.iov_len = virtio_net_hdr length`
2. `out_sg.iov_base = HVA skb->data`
`out_sg.iov_len = skb->len`

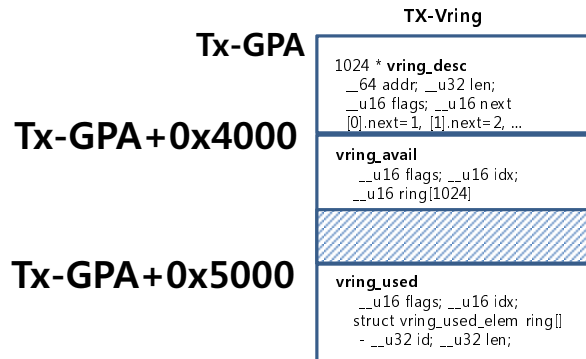


qemu_sendv_packet_async
(..., out_sg, out_len, ...
virtio_net_tx_coplete)

Virtio – moving data

Guest – convert GVA -> HPA

Host GPA -> HVA, HVA -> GPA



```
VRing vring;
{
    unsigned int num=...
    hwaddr desc = Desc Tx-GPA
    hwaddr avail = Desc Tx-GPA + ofst
    hwaddr used = Tx-GPA + ofst
} /* VRing */
```

Host Vring Operations

```
virtio_net_flush_tx(...)
virtqueue_pop(q->tx_vq, &elem)
hwaddr desc_pa = vq->vring.desc;
i = virtqueue_get_head(vq, vq->last_avail_index++)
    - hwaddr pa = vq->vring.avail + offsetof(VRingAvail, ring[i])
    - return lduw_phys(pa)
hwaddr desc_pa = vq->vring.desc
```

Convert to GVA

```
flags = vring_desc_flags(desc_pa, i)
pa = desc_pa + sizeof(VRingDesc) * i + offsetof(VringDesc, flags)
return lduw_phys(pa)
```

Convert GPA – &vring_desc->addr to GVA

```
elem->out_addr[elem->out_num] = vring_desc_addr(desc_pa, i)
elem->out_sg[...].iov_len = vring_desc_len(desc_pa, i)
```

Convert GPA – vring_desc->addr to GVA

```
elem->out_sg[...].iov_base = cpu_physical_memory_map(elem->out_addr[...], ...)
```

-
- Tx out Backend
- Notify guest – Tx interrupt completion

Guest Vring Operations

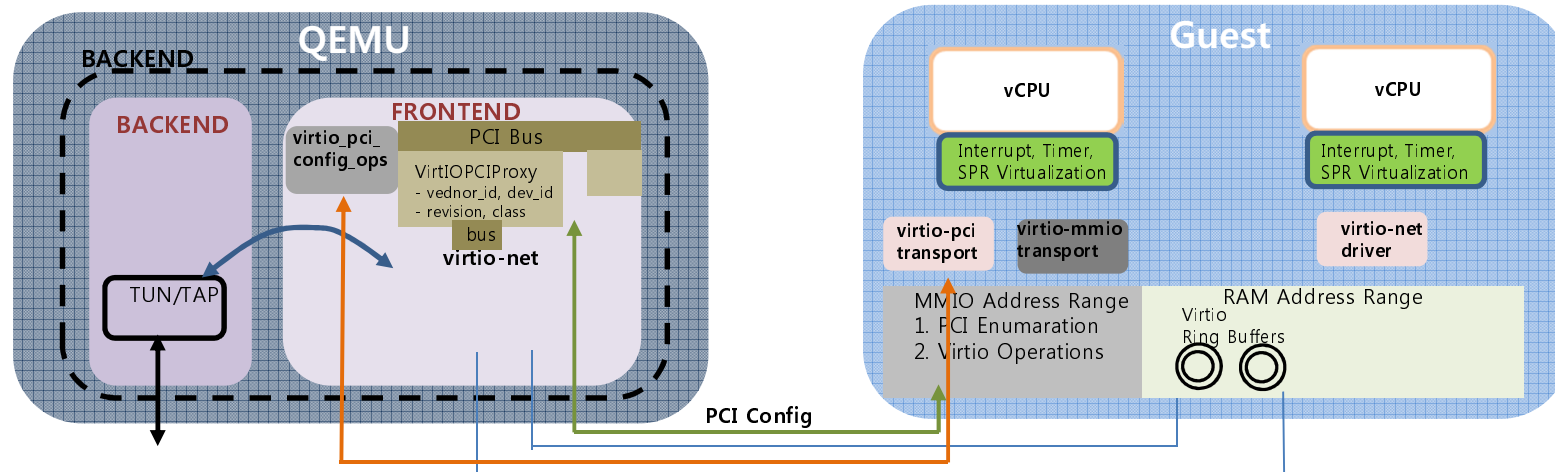
```
xmit_skb(...)
sg_set_buf(scatterlist *sg, ..., virtio_net_hdr)
    - sg->page_link = page
    - sg->offset = page offset
    - sg->length ...
sg_set_buf(scatterlist *sg, ..., skb->data)
.....
vq->vring.desc[i].flags = ...
vq->vring.desc[i].addr = GPA of page
vq->vring.desc[i].addr = sg->length
....
vq->notify(vq)
    - mmio write – VIRTIO_XXX_QUEUE_NOTIFY
```

Virtio and Device Pass-through

Basic Operation	<ul style="list-style-type: none"> - Backend/Guest direct access to shared Vring buffers - PIO - Switching at software level - Management Flexibility – internal SDN support <code>ovs-vsctl add-port br0 <phys-intfc> - vSwitch</code> <code>ovs-ofctl - control flows</code> - IRQ bottleneck – QEMU – call into kvm inject Kernel – inject directly 	<ul style="list-style-type: none"> - Direct access to hw memory regions - DMA Support - Switching at hw level – SR-IOV depends on #of Queues - Management Flexibility – external SDN capable - IRQ bottleneck – hw enhancements, posted interrupts, exitless EOI improve things – closer to native
Migration	<ul style="list-style-type: none"> - Virtio lockless - Saves device state, tracks dirty pages 	<ul style="list-style-type: none"> - QEMU sets 'unmigratable', or installs migration blocker - Guest can be holding a lock – deadlock, hw state,
Scalability	<ul style="list-style-type: none"> - Practical limitations – primarily performance 	<ul style="list-style-type: none"> - Number of Devices limited, limits #VMs - SR-IOV - #of VF - # of queues
Network Performance	<ul style="list-style-type: none"> - Soft switching – bridge, vSwitch - Several IO HOPS - Can approach near native – 10Ge for few bridged Guest 	<ul style="list-style-type: none"> - Switching done at HW level – hw queues - Performance scales with # of Guests - DMA support - IRQ Passthrough still a problem
Host Performance	<ul style="list-style-type: none"> - PIO – takes cpu cycles - Exits – few but still - Guest pages swappable 	<ul style="list-style-type: none"> - Guest pinned – can't swap - Fewer exits - Less PIO
Cloud Environment	<ul style="list-style-type: none"> - Cloud friendly – migration, SDN, paging 	<ul style="list-style-type: none"> - Not Cloud friendly, great for NFV/RT DPK, run to completion

Virtio PCI Architecture

- virtio-net example with QEMU backend – virtio-pci

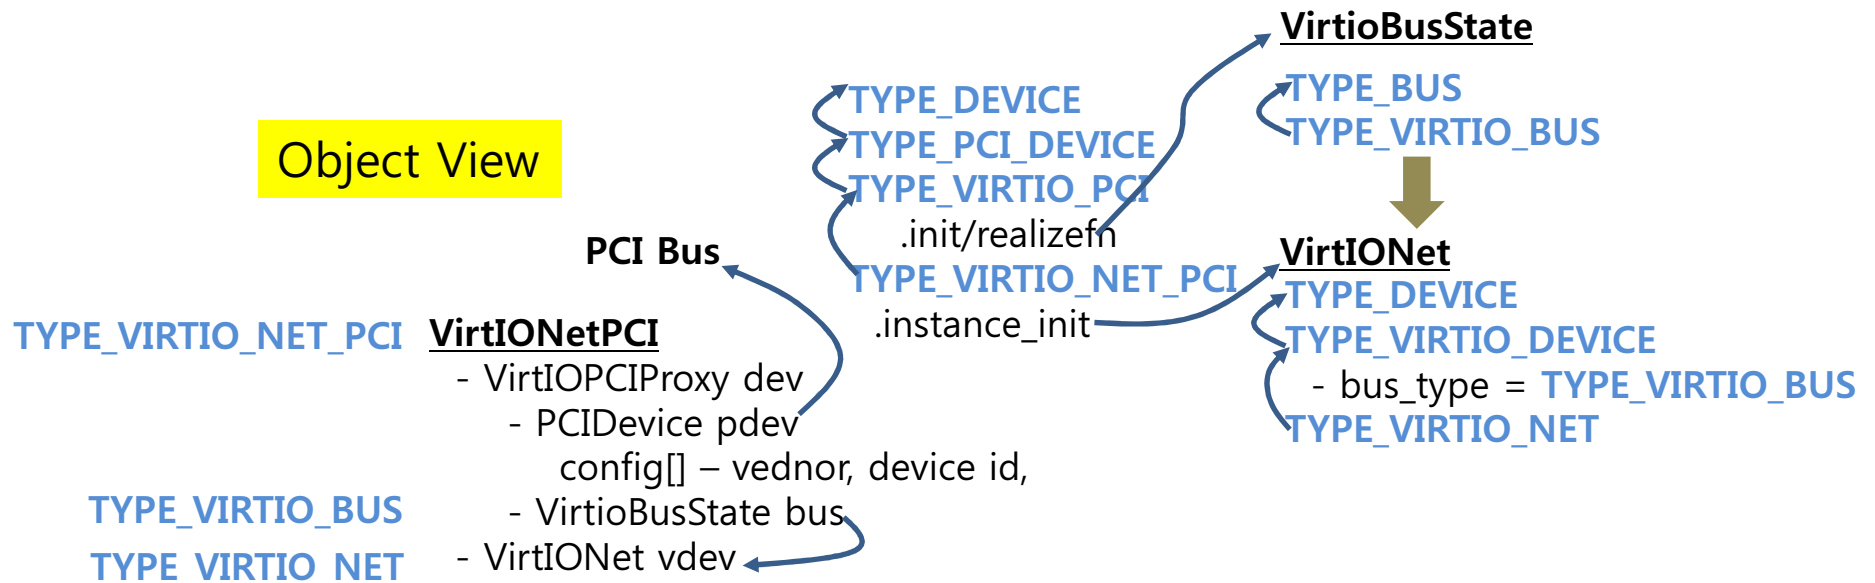


- Virtio device – combination of mmio & paravirt device
- Before Guest Runs QEMU does
 - creates proxy that plugs into PCI Bus
 - During instantiation of VirtIOProxy its
 - PCIDevice vendor id, device id, class, ... are set
 - Instantiates virtio-net – bus_type = TYPE_VIRTIO_BUS
 - Plugs into VirtIOProxy bus – TYPE_VIRTIO_BUS
 - Fills in PCI BAR0 type PIO
 - Associates virtio_pci_config_ops with B/D/F BAR0
- Guest
 - Enumerates PCI Bus – discovers virtio-net – via mmio
 - Loads virtio-pci, creates virtio-net device
 - virtio-net driver loads probes virtio-net backend – via mmio

QEMU Object Model

- QEMU Class, Object view of '-device virtio-net-pci'
 - First instantiate Class – C++ Class definition
 - Next the Object – C++ Declare Class variable
 - Realize it– C++ constructor default or defined

Class View

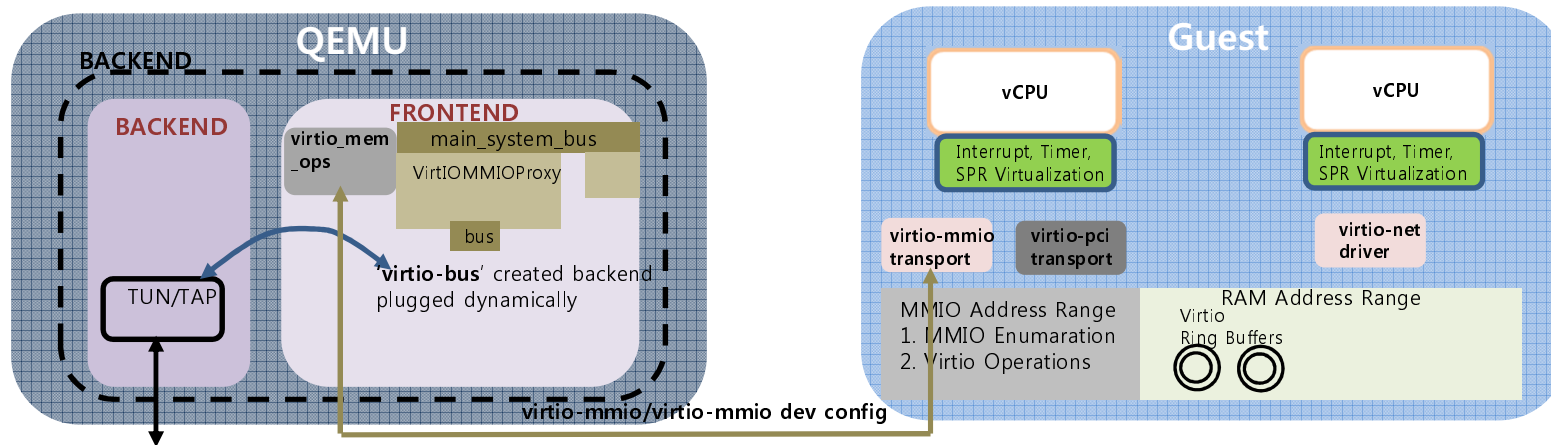


virtio-mmio transport


- virtio-net example with QEMU backend – virtio-mmio
- Discovery/Probing ... like PCI
- Primarily ARM – with Guest QEMU/Guest PCI support – virtio-mmio less use
- Some Use cases
 - ❑ Want your own Machine Model – don't want PCI, have Device Tree support
 - ❑ Lots of Embedded Devices – simplified machine model
 - Automotive, Edge Network, Set top Box
 - ❑ virtio-mmio another option

Virtio MMIO Architecture

- Virtio-mmio – example similar to PCI

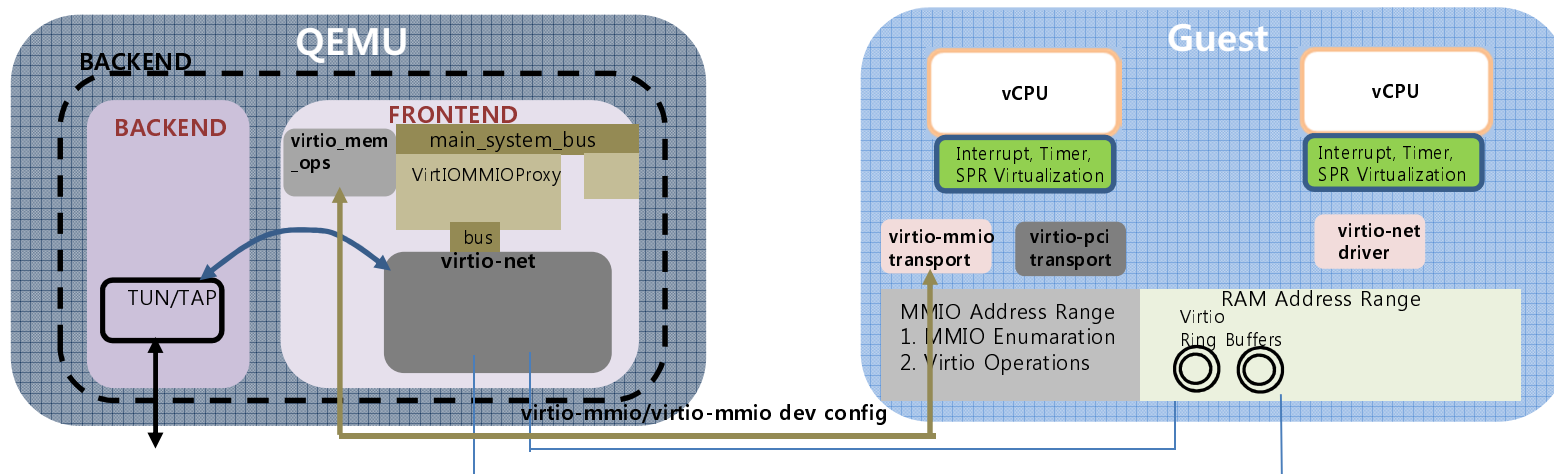


1. Instantiate multiple virtio-mmio devices – no qemu args implicitly done




2. Associate resources with each virtio-mmio range
 - MMIO address range a page, interrupt # - only machine models what resources

Virtio MMIO Architecture



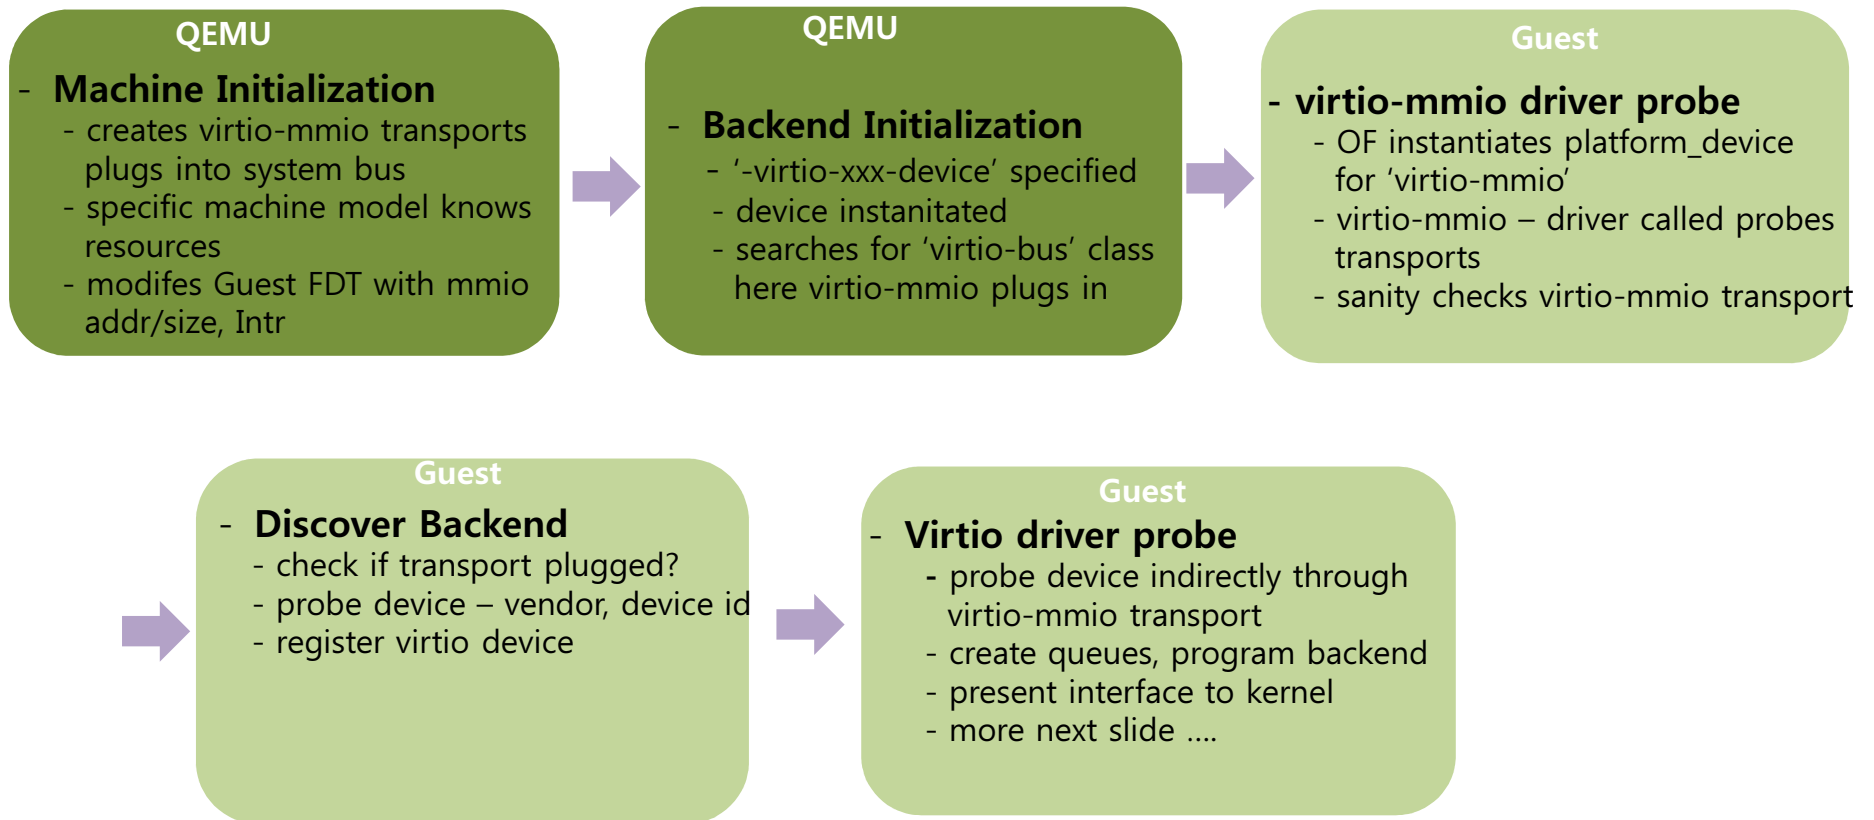
3. '-virtio-net-device' – instantiates/plugs TYPE_VIRTIO_NET

- No transport specified any backend (virtio-net, virtio-blk,...) plug into transport
- Virtio-net inherits VirtIODevice which sets 'bus_type = TYPE_VIRTIO_BUS'
- Finds matching bus VirtIOMMIOproxy->bus, plugs TYPE_VIRTIO_NET
- Finds and binds to QEMU backend – f.e. –netdev type=tap ...



Guest virtio discovery framework – virtio-mmio view

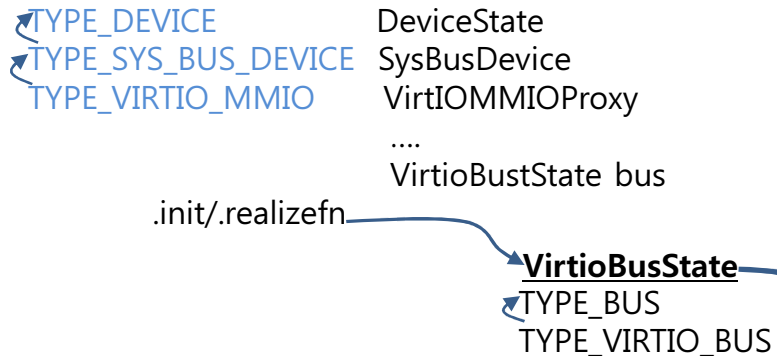
- Transparent to Guest – enable virtio and mmio
- Device Tree used



Guest virtio discovery framework – virtio-mmio view

Machine Initialization

```
....  
sysbus_create_simple("virtio-mmio", base, pic[irq])  
....
```




```
add_virtio_mmio_node(fdt, ..., mmio addr, irq pin)
```

Backend Initialization

```
.....  
QEMU option ... '-virtio-net-device'
```

```
device_init_func(opts, ....)
```



virtio-mmio driver probe

```
of_platform_populate(..., of_device_id match[], ...)
```

```
DT Node  
virtio-mmio {  
    addr, size, irq;  
}  
platform_device {  
    ....  
}
```

```
... virtio_mmio_driver = {  
    .probe = virtio_mmio_probe,  
    ...  
}
```

```
virtio_mmio_probe(*pdef)  
- virtio_mmio_vdev *vm_dev  
vm_dev->base = ioremap(virtio-mmio - GPA, size)  
virtio_device transport interface, PCI has one too  
vm_dev->vdev.config = &virtio_mmio_config_ops  
sanity check – mmio to 'virtio_mem_ops' handlers  
magic = readl(vm_dev->base + VIRTIO_MMIO_MAGIC_VALUE)  
version = readl (.....)
```

Guest virtio discovery framework – virtio-mmio view

Discover Backend

Identify if device plugged, if yes identify device

```
vm_dev->vdev.id.device = readl(vm_dev->base + VIRTIO_MMIO_DEVICE_ID)
vm_dev->vdev.id.vendor = readl(...)
```

register the device

```
register_virtio_device(struct virtio_device dev=vm_dev->vdev)
```

Ack device found by transport, use transport interface

```
- dev->config->set_status( ... get_status() | VIRTIO_CONFIG_S_ACKNOWLEDGE)
```

find matching driver on virtio bus

- bus_for_each_drv(...)
- virtio_dev_match(dev, drv)

Ack driver found for device

```
- dev->config->set_status(...get_status() | VIRTIO_CONFIG_S_DRIVER)
```

Feature Negotiation – these are key performance features

- Get backed features – be_features = dev->config->get_features(vdev)
- walk driver feature table - check if backend supports – be_features bit set
 - if supported set vdev->features
- select features – vdev->config->finalize_features(vdev->features[])
 - a) backend features not supported by driver don't get selected
 - b) driver features not supported by backend don't get selected
- call driver probe – virtnet_probe()

virtio driver probe

instantiate network device interface

```
dev = alloc_etherdev_mq(..., # of queues)
```

....

Various performance features – primarily offload, big packets

- Check supported features – from vdev->features – set dev->hw_features
- Vdev->config->find_vqs(...)
 - Initialize queues – allocated by guest
 - Tell backend GFN of Vring and buffer count for each queue
 - Backend – sets GPA and GPA indexes into Descriptors, Available, Used ring.

Virtio Performance

- When transport/backend are not 'fused' performance features not exported
 - ❑ Due to way QEMU instantiates objects – properties set at TYPE_DEVICE class
 - ❑ After device plugged – properties not set
 - ❑ If transport/backend not fused – properties/performance features not used
 - ❑ Created patch for virtio-mmio – applies when backend plugged
 - https://github.com/mjsmar/virtio_net_fix.git

Virtio Performance

- Performance features
 - ❑ Red Hat multi-queue tap
 - tap arg – 'queues=n' for scalability
 - Creates multiple queue virtio/tap tx/rx queue pairs
 - vCPU scaling for tx/rx, serializes flows - TCP sessions, UDP connections
 - ❑ tx=timer,x-txtimer=<n>uS
 - Host kicks the backend periodically – limit exits
 - you can adjust how often backend polls tx virtqueue – tune latency vs. CPU
 - ❑ Offload – bigger pkts few exits, offload to host
 - Ring descriptors have – 1 for virtio_net_hdr other for data
 - probe tun/tap for vnet_hdr support – for offloads - IFF_VNET_HDR
 - Probe tun/tap for GSO – TCP,UDP, - TSO, UFO
 - Options eventually make it to virtio-net net_device 'features'
 - virtio_net_hdr – flags for CSUM & define range
 - check skb fragments for for GSO – set vnet_hdr_net gso_type, size

Reconnaissance of Virtio: What's new and how it's all connected?



Q & A

Thank you.

Mario Smarduch
Senior Virtualization Architect
Open Source Group
Samsung Research America (Silicon Valley)
m.smarduch@samsung.com